

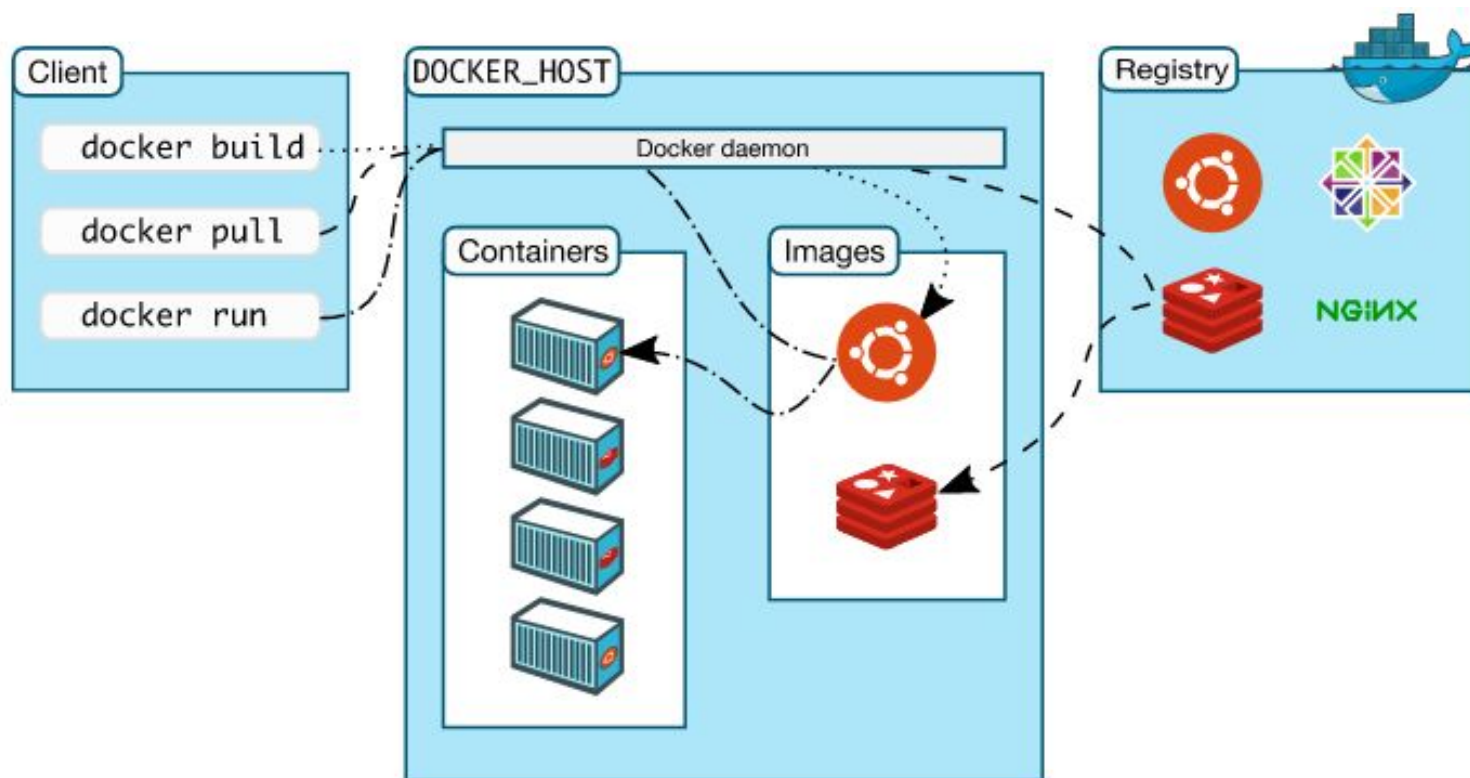
Kubernetes

Solar Team

Prepare Your Environments

- Windows/MacOS, Ram 16
- Docker Desktop + enable Kubernetes
- The Kubernetes command-line tool, [kubectl](#).
- VPN Connection

Docker Architecture



- Containers are a good way to bundle and run your applications.
- How can we deal with large numbers of container and distributed issues?

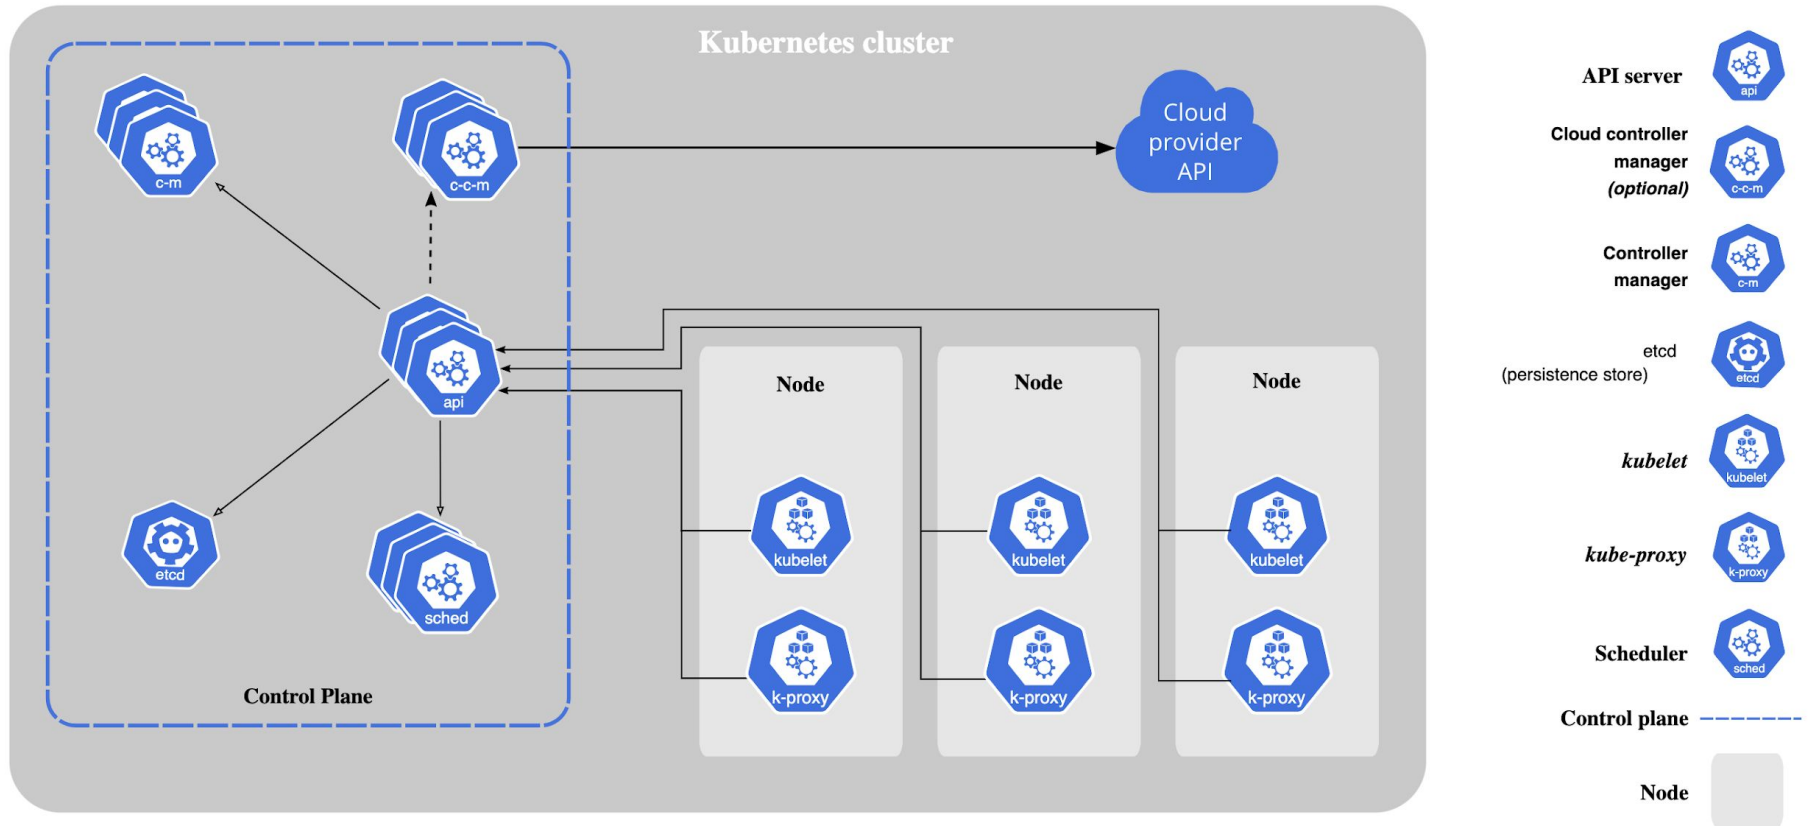
Kubernetes Overview

- An **open source container orchestration platform** designed for running distributed applications and services at scale.
- The core functionality is scheduling workloads in containers across your infrastructure.
- Google open-sourced the Kubernetes project in 2014.
- Known as K8s.
- Latest version is v1.20 (February 2021)
- <https://kubernetes.io/>

Why do we need Kubernetes?

- Kubernetes provides you with:
 - **Service discovery and load balancing**
 - **Storage orchestration**
 - **Automated rollouts and rollbacks**
 - **Automatic bin packing**
 - **Self-healing**
 - **Secret and configuration management**

Kubernetes Architecture



kubectl

- **kubectl** is the command-line tool to interact with the Kubernetes clusters.
- It makes the necessary Kubernetes API calls for you.
- `kubectl [command] [TYPE] [NAME] [flags]`
- Example:
 - `kubectl version --short`
 - `kubectlt get pods`
 - `kubectl get po`
 - `kubectl get po etcd-docker-desktop --namespace kube-system`
 - `kubectl get svc #(or service or services)`
 - `kubectl get po --namespace kube-system`
 - `kubectl describe namespaces default`
 - `kubectl cluster-info --kubeconfig ~/.kube/config.180`
 - `kubectl run nginx --image=nginx --replicas=3 --labels=app=web-proxy --dry-run -o=yaml`

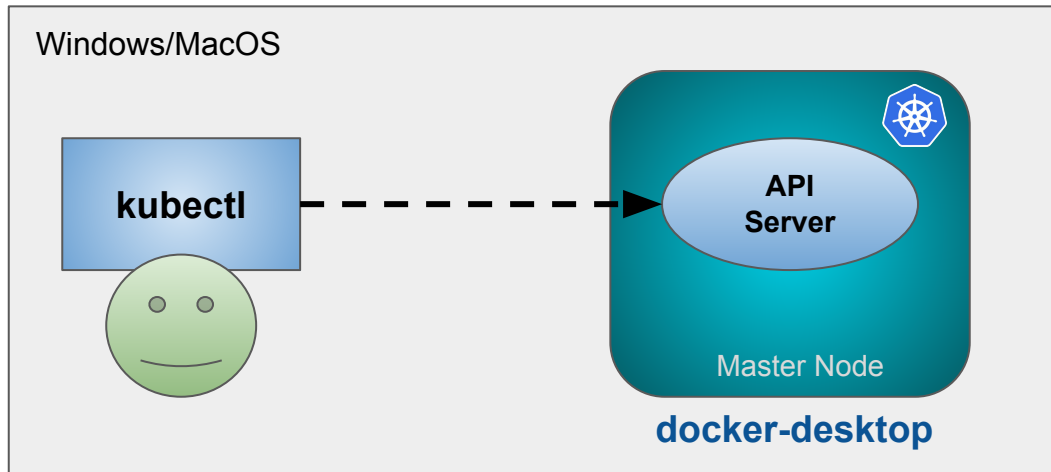
Kubernetes Dashboard

- **For kube 1.14**
 - `kubectl apply -f`
<https://raw.githubusercontent.com/kubernetes/dashboard/v1.10.1/src/deploy/recommended/kubernetes-dashboard.yaml>
- **Access Dashboard**
 - `kubectl proxy`
 - <http://localhost:8001/api/v1/namespaces/kube-system/services/https:kubernetes-dashboard:/proxy/>
- **Get Token**
 - `kubectl -n kube-system describe secret $(kubectl -n kube-system get secret | grep admin-user | awk '{print $1}')`
- **Another tool is [Lens](#) (Kubernetes IDE).**

Kubernetes API

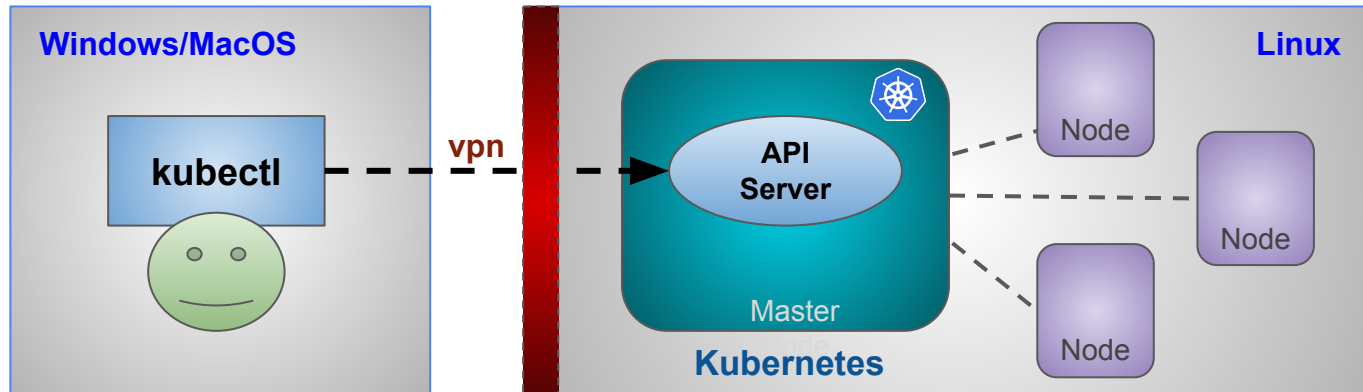
- The core of Kubernetes' control plane is the **API server**.
- The API server exposes an HTTP API.
- Used for communication between user (query and manipulate the kubernetes objects), internal components and external components.
- Most operations can be performed through the **kubectl** and other tool like **kubeadm**.
- Can also access the API directly using REST calls.
- Provided client library for several programming languages like Go, Java, Python, etc.
- Kubernetes stores the serialized state of objects by writing them into **etcd**.
- Support both response format, json and protobuf.

Access Local Kubernetes Cluster



- Start docker and enable local kubernetes.
- Open a terminal console then try these commands.
 - `kubectl cluster-info` # Print the address of the master and cluster services
 - `kubectl config current-context` # view the current context
 - `kubectl config view` # Display merged kubeconfig settings or a specified kubeconfig file
- To access api server.
 - `kubectl proxy` # default port is 8001
 - Open this url in browser - <http://localhost:8001/api/v1/>
 - `kubectl proxy --port=8080` # change proxy port to 8080

Access Remote Kubernetes Cluster



- Access a remote cluster with Kubectl
 - Copy Kubernetes config file from the master node to for example **(user-home)/.kube/config.180**
 - **Solution #1** - replace local config file with remote cluster config file.
 - **Solution #2** - setting **KUBECONFIG** environment variable.
 - `export KUBECONFIG=~/.kube/config.180`
 - **Solution #3** - using `--kubeconfig` option.
 - `kubectl cluster-info --kubeconfig ~/.kube/config.180`
 - **Solution #4** - using proxy (access kubernetes api)
 - `kubectl proxy`
 - Or using `kubectl proxy --kubeconfig ~/.kube/config.180`

Use with caution!

Resource Short-names

Short Name	Full Name	Short Name	Full Name
csr	certificatesigningrequests	pvc	persistentvolumeclaims
cs	componentstatuses	pv	persistentvolumes
cm	configmaps	po	Pods
ds	daemonsets	pdb	poddisruptionbudgets
deploy	deployments	psp	podsecuritypolicies
ep	endpoints	rs	replicasets
ev	events	rc	replicationcontrollers
hpa	horizontalpodautoscalers	quota	resourcequotas
ing	ingresses	sa	serviceaccounts
limits	limitranges	svc	services
ns	namespaces		
no	nodes		

Basic Commands

- **kubectl version --short** - Print the client and server version information for the current context
- **kubectl options** - Print a list of global command-line options (applies to all commands).
- **kubectl api-versions** - Print the supported **API versions on the server**, in the form of "group/version".
- **kubectl api-resources** - Print the supported API resources on the server.
- **kubectl proxy** - Creates a proxy server or application-level gateway between localhost and the Kubernetes API Server.
- **kubectl explain** - List the fields for supported resources.

Kubernetes Release Notes

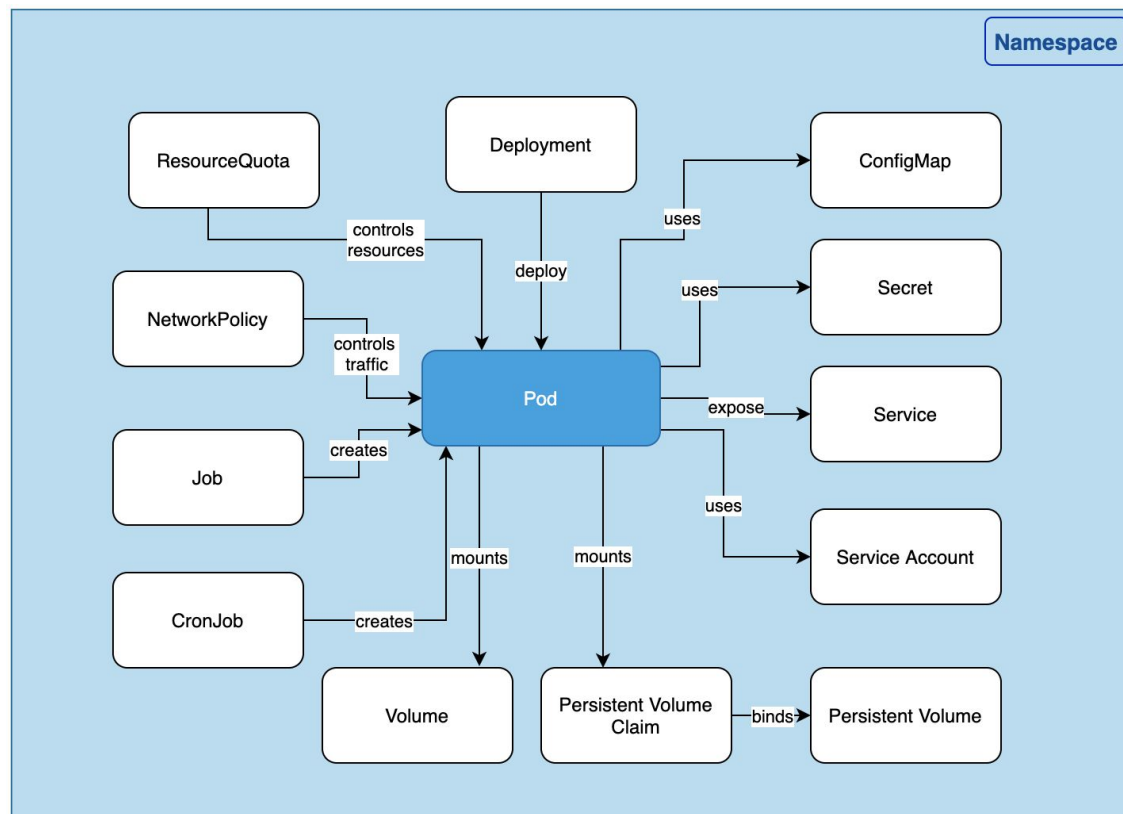
- <https://v1-18.docs.kubernetes.io/docs/setup/release/notes/>
- kube-apiserver no longer serves the following deprecated APIs: (#85903, @liggitt) * All resources under apps/v1beta1 and apps/v1beta2 - use apps/v1 instead * daemonsets , deployments , replicaset resources under extensions/v1beta1 - use apps/v1 instead * networkpolicies resources under extensions/v1beta1 - use networking.k8s.io/v1 instead * podsecuritypolicies resources under extensions/v1beta1 - use policy/v1beta1 instead

kubeconfig And Context

- **Kubeconfig files** keep an organize information about clusters, users, namespaces, and authentication mechanisms.
- The kubectl command-line tool uses kubeconfig files to find the information it needs to choose a cluster and communicate with the API server of a cluster.
- The kubectl looks for a file named config in the `$HOME/.kube` directory.
- Specify other kubeconfig files by setting the **KUBECONFIG** environment variable or by setting the **--kubeconfig** flag.
- **Context** element in a kubeconfig file is used to group access parameters under a convenient name.
- Each context has three parameters: **cluster, namespace, and user**.
- By default, the kubectl command-line tool uses parameters from the current context to communicate with the cluster.

Kubernetes Primitives(Resources)

- The basic building blocks in the Kubernetes architecture for creating and operating an application on the platform.



Some of Kubernetes Primitives

Kubernetes Object

Structure of Kubernetes Object

apiVersion:

apps/v1, v1, ...

kind:

deployment, pod, service, .

metadata:

name, namespace, labels,...

spec: { Desired State }

selector, replicas, template,...

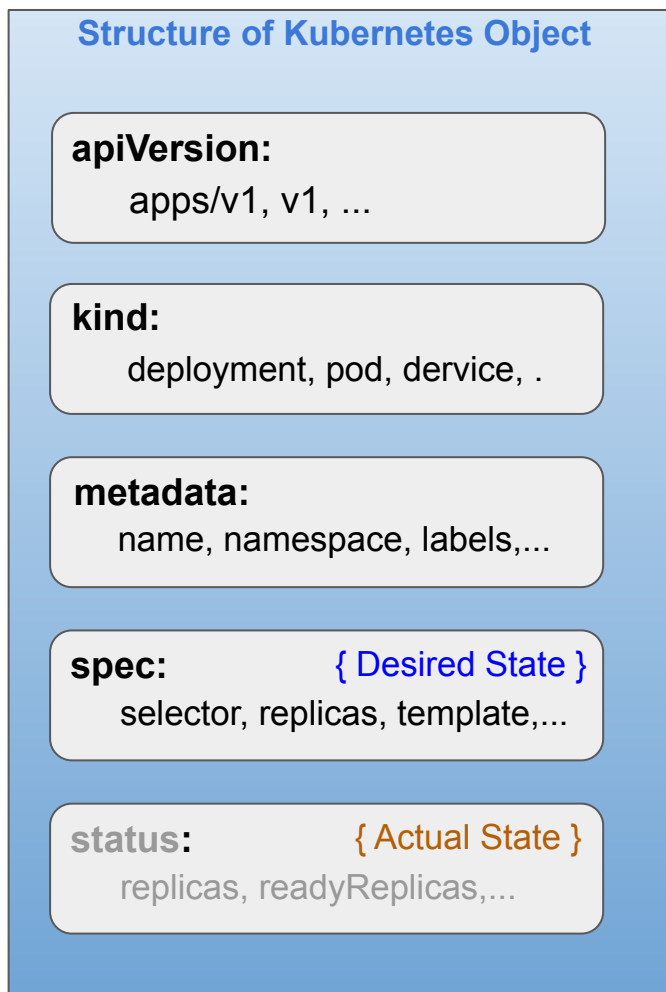
status: { Actual State }

replicas, readyReplicas,...

- Kubernetes objects are **persistent entities** in the Kubernetes system.
- Kubernetes uses these entities to **represent the state** of your cluster.
- Kubernetes uses the **YAML file** for declaring.
- A Kubernetes object is a "**record of intent**"
- Use the **kubectl** command-line to work(create, modify, delete, etc.) with Kubernetes objects.

<https://v1-18.docs.kubernetes.io/docs/setup/release/notes/>

Kubernetes Object Structure



- **API version** - Which version of the Kubernetes API you're using to create this object.
- **Kind** - What kind of object you want to create e.g., a Pod or a Service.
- **Metadata** - Data that helps uniquely identify the object, including a **name** string, **UID**, and optional namespace.
- **The specification ("spec" for short)** - What state you desire for the object. Which image should run in the container, or which environment variables should be set for?
- **Status** - Describes the actual state of an object.

Kubernetes Object Management

- Using **kubectl** command-line tool supports several different ways to create and manage Kubernetes objects.
- Two approach - **Imperative approach** and **Declarative approach**.
- **Warning:** Use only one technique, mixing techniques for the same object results in undefined behavior.

Imperative approach

- Two kinds of Imperative approaches, **Imperative commands** and **Imperative object configuration**
- **Imperative commands** does not require a manifest definition(YAML file).
 - Use the `kubectl run` or `kubectl create` command to create an object on the fly.
 - `kubectl create deployment nginx --image nginx`
 - `kubectl run web --image=nginx --restart=Never --port=80 # run-web_get.yml`
- **Imperative object configuration** creates objects from a manifest file (YAML file)
 - `kubectl create -f nginx.yaml`
 - `kubectl delete -f nginx.yaml -f redis.yaml`
 - `kubectl replace -f nginx.yaml` # Update the objects defined in a configuration file by overwriting the live configuration:
- **The recommended way for development environments.**
- **Dry run.** Print the corresponding API objects without creating them.
 - `kubectl create nginx --image=nginx --dry-run`

Declarative approach

- **Declarative object configuration** - create, update, and delete operations are automatically detected per-object by kubectl.
 - `kubectl apply -f configs/`
 - `kubectl diff -f configs/`
 - `kubectl apply -R -f configs/` # Recursively process directories.
 - **This approach is a recommended way for production environments.**

Namespace

Namespaces

- **Namespace** is a **virtual cluster concept** in Kubernetes.
- Kubernetes supports multiple virtual clusters backed by the same physical cluster.
- Provide **isolated resources that are spread across teams and projects** - via resource quotas, network policies.
- Names of resources need to be **unique within a namespace**, but not across namespaces.
- Four initial namespaces are **default**, **kube-system**, **kube-public** and **kube-node-lease**.
- **Not all objects are in a namespace** for example Node, PersistentVolume and Namespace itself.
- Prefix “**kube-**” is reserved for Kubernetes system namespaces.

Namespace Commands


- Viewing namespaces
 - `kubectl get namespace`
- Setting the namespace for a request
 - `kubectl run nginx --image=nginx --namespace=<insert-namespace-name-here>`
 - `kubectl get pods --namespace=<insert-namespace-name-here>`
- Setting the namespace preference
 - `kubectl config set-context --current --namespace=<insert-namespace-name-here>`
 - `kubectl config view --minify | grep namespace # verify it`
- List objects in a namespace
 - `kubectl api-resources --namespaced=true`
- List objects not in a namespace
 - `kubectl api-resources --namespaced=false`

Create A Namespace Example

```
kubectl create namespace dev --dry-run -o=yaml > namespace-dev.yml
```

```
apiVersion: v1
kind: Namespace
metadata:
  creationTimestamp: null
  name: dev
spec: {}
status: {}
```

namespace-dev.yml



```
kubectl get ns dev -o yaml > namespace-dev_get.yml
```

```
apiVersion: v1
kind: Namespace
metadata:
  annotations:
    kubectl.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"v1","kind":"Namespace",
       "metadata":{"annotations":{},"name":"dev"}}
  creationTimestamp: "2021-02-28T08:31:45Z"
  name: dev
  resourceVersion: "638999"
  selfLink: /api/v1/namespaces/dev
  uid: 641f333d-799f-11eb-bd28-025000000001
spec:
  finalizers:
    - kubernetes
status:
  phase: Active
```

namespace-dev_get.yml

Deployment

Deployment

- A Deployment is responsible for creating and updating instances of your application (Pods).
- Deployment adds several important features for a production environment.
 - Scalability
 - Update strategy
 - Historical
- Deployment creates:
 - A **ReplicaSet** to ensure that the desired number of Pods is always available,
 - Specifies a strategy to replace Pods (such as RollingUpdate).



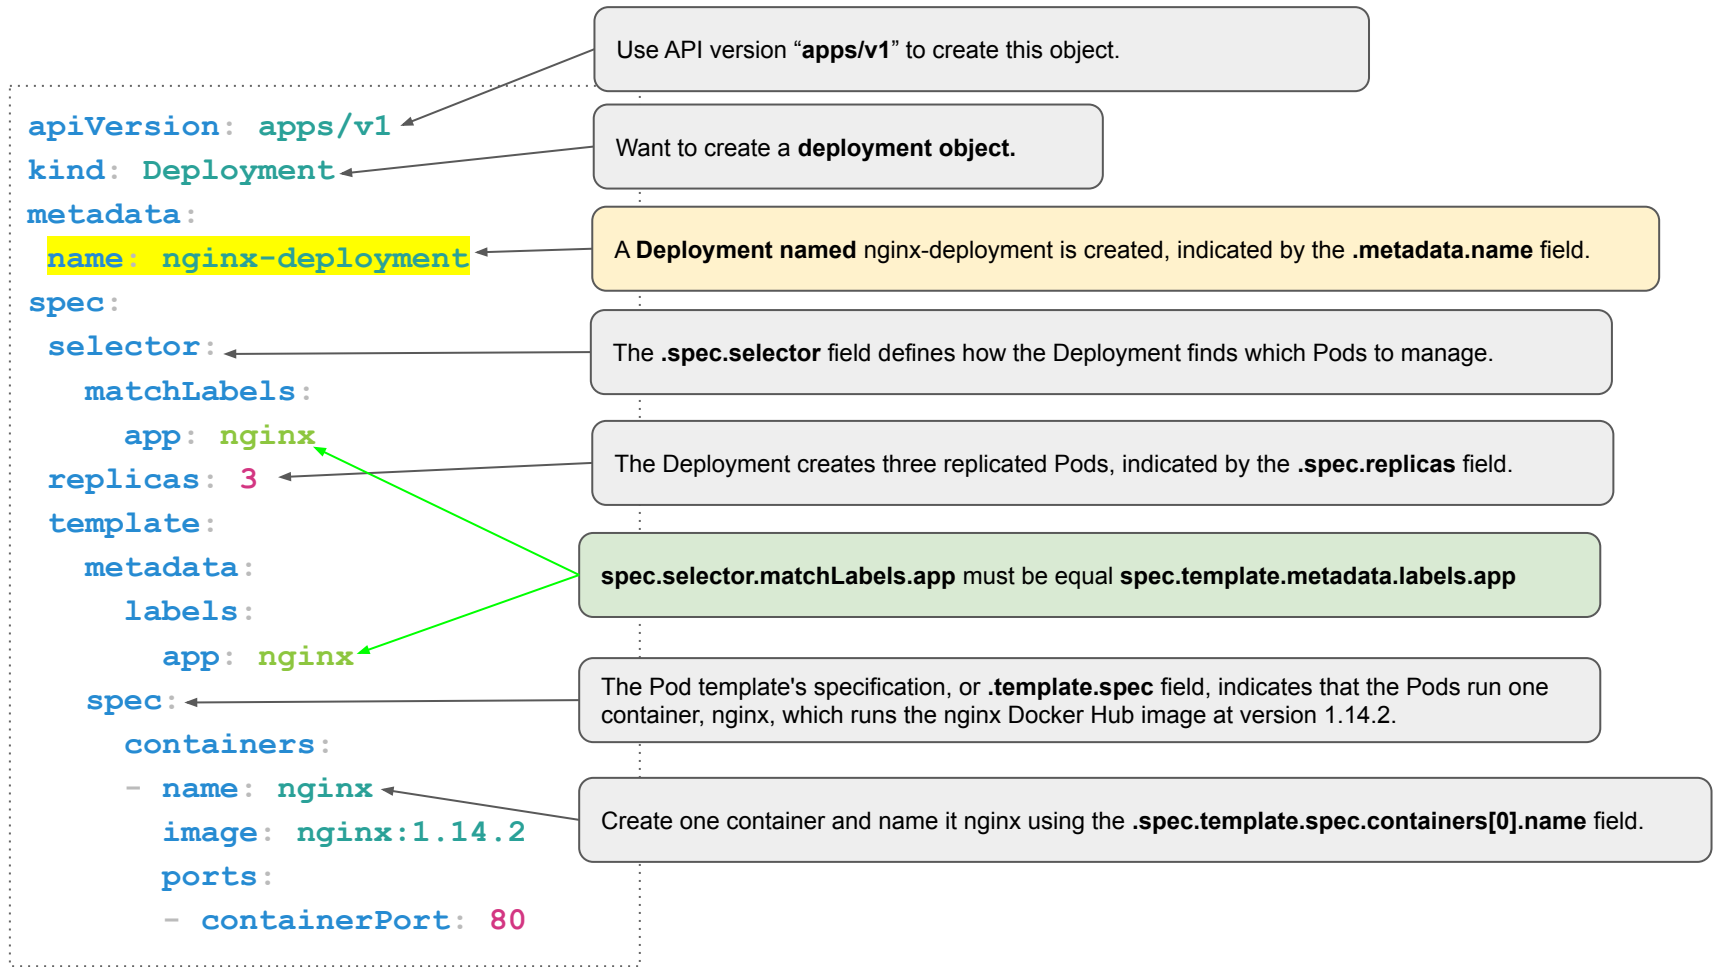
Create Deployment(Imperative Way)

- `kubectl create deployment web-server --image=nginx`
- `kubectl get deployment web-server -o yaml > web-server_get.yml`

```
1  apiVersion: extensions/v1beta1
2  kind: Deployment
3  metadata:
4    annotations:
5      deployment.kubernetes.io/revision: "1"
6    creationTimestamp: "2021-03-02T13:48:40Z"
7    generation: 1
8    labels:
9      app: web-server
10   name: web-server
11   namespace: default
12   resourceVersion: "752957"
13   selfLink: /apis/extensions/v1beta1/namespaces/default/deployments/web-server
14   uid: fee753c6-7b5d-11eb-bd28-025000000001
15  spec:
16    progressDeadlineSeconds: 600
17    replicas: 1
18    revisionHistoryLimit: 10
19    selector:
20      matchLabels:
21        app: web-server
22    strategy:
23      rollingUpdate:
24        maxSurge: 25%
25        maxUnavailable: 25%
26      type: RollingUpdate
27    template:
28      metadata:
29        creationTimestamp: null
30        labels:
31          app: web-server
32      spec:
33        containers:
34          - image: nginx
35            imagePullPolicy: Always
36            name: nginx
37            resources: {}
38            terminationMessagePath: /dev/termination-log
39            terminationMessagePolicy: File
40          dnsPolicy: ClusterFirst
41          restartPolicy: Always
42          schedulerName: default-scheduler
43          securityContext: {}
44          terminationGracePeriodSeconds: 30
45  status:
46    availableReplicas: 1
47    conditions:
48    observedGeneration: 1
49    readyReplicas: 1
50    replicas: 1
51    updatedReplicas: 1
```

web-server_get.yml

Basic Deployment Manifest



nginx-deployment.yml


Create Deployment from YAML (Declarative Way)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

nginx-deployment.yml



kubectl create -f nginx-deployment.yml



```
✓ kubectl get deployment
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deployment    3/3     3             3           13m
```

```
✓ kubectl get rs
NAME                                DESIRED   CURRENT   READY   AGE
nginx-deployment-756d9fd5f9         3         3         3       4h10m
```

```
✓ kubectl get po
NAME                                READY   STATUS    RESTARTS   AGE
nginx-deployment-756d9fd5f9-2fkx6   1/1     Running   0          13m
nginx-deployment-756d9fd5f9-9pl8x   1/1     Running   0          13m
nginx-deployment-756d9fd5f9-vjq6p   1/1     Running   0          33s
```

Compare Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.14.2
        ports:
        - containerPort: 80
```

nginx-deployment.yml

kubectl get deployment nginx-deployment
-o=yaml > nginx-deployment_get.yml

```
1  apiVersion: extensions/v1beta1
2  kind: Deployment
3  metadata:
4    annotations:
5      deployment.kubernetes.io/revision: "2"
6      kubectl.kubernetes.io/last-applied-configuration: |
7        {"apiVersion":"apps/v1","kind":"Deployment","metadata":{"annotations":{},"name":
8  creationTimestamp: "2021-03-01T10:08:58Z"
9  generation: 3
10 labels:
11   app: nginx
12   name: nginx-deployment
13   namespace: default
14   resourceVersion: "693345"
15   selfLink: /apis/extensions/v1beta1/namespaces/default/deployments/nginx-deployment
16   uid: 23a62fa5-7a76-11eb-bd28-025000000001
17 spec:
18   progressDeadlineSeconds: 600
19   replicas: 3
20   revisionHistoryLimit: 10
21 > selector: ...
24 strategy:
25   rollingUpdate:
26     maxSurge: 25%
27     maxUnavailable: 25%
28   type: RollingUpdate
29 template:
30   metadata:
31     creationTimestamp: null
32     labels:
33       app: nginx
34 > spec: ...
50 status:
51   availableReplicas: 3
52   conditions:
53 > - lastTransitionTime: "2021-03-01T10:22:24Z" ...
59 > - lastTransitionTime: "2021-03-01T10:08:58Z" ...
65   observedGeneration: 3
66   readyReplicas: 3
67   replicas: 3
68   updatedReplicas: 3
```

ReplicaSet

- Replicate a **guaranteed number of Pods** with the same configuration.
- Deployment is the higher-level concept that manages the ReplicaSet internally with no involvement required by the end user.
- ReplicaSets provide self-healing capabilities.



Lab: Deployment

- `kubectl create -f nginx-deployment.yml` Or `kubectl apply -f nginx-deployment.yml`
- `kubectl get deployments,pods,replicasets`
- `kubectl describe deployment.apps/nginx-deployment`
- `kubectl set image deployment nginx nginx=nginx:1.19.2`
- `kubectl rollout history deployment nginx`
- `kubectl get deployments,pods,replicasets`
- `kubectl get pods -A`
- `kubectl describe deployment nginx`
- `kubectl describe deployment.app/nginx-deployment`
- `kubectl rollout history deployments nginx-deployment --revision=2`
- `kubectl port-forward <POD> 8002:80`
- `kubectl rollout undo deployment my-deploy --to-revision=1`

nginx-deployment.yml

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx
  annotations:
    app-domain: "infra"
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 2
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

Names and UIDs

- Each object in your cluster has a **Name** that is unique for that type of resource.
- Every Kubernetes object also has a **UID**(systems-generated uuid string) that is unique across your whole cluster.
- For example, you can only have one Pod named myapp-1234 within the same [namespace](#), but you can have one Pod and one Deployment that are each named myapp-1234.
- For non-unique user-provided attributes, Kubernetes provides **labels** and **annotations**.
- A client-provided string that refers to an object in a resource URL, such as /api/v1/pods/**some-name**.

Label

- Labels are **key/value pairs** that are attached to objects, such as deployment.
- Label **key must be unique** for a given object.
- Labels can be attached to objects at creation time and subsequently added and modified at any time.
- Can manually set and automatically set by Kubernetes itself.
- Used to organize and to locate subsets of objects.
- ReplicaSet monitor the pod by using labels.

```
"metadata": {  
  "labels": {  
    "key1" : "value1",  
    "key2" : "value2"  
  }  
}
```

Recommended Labels

- There are recommended labels. They make it easier to manage applications but aren't required.
- Shared labels and annotations share a common prefix: **app.kubernetes.io**.
- Labels without a prefix are private to users.

Key	Description	Example	Type
<code>app.kubernetes.io/name</code>	The name of the application	<code>mysql</code>	string
<code>app.kubernetes.io/instance</code>	A unique name identifying the instance of an application	<code>mysql-abcxzy</code>	string
<code>app.kubernetes.io/version</code>	The current version of the application (e.g., a semantic version, revision hash, etc.)	<code>5.7.21</code>	string
<code>app.kubernetes.io/component</code>	The component within the architecture	<code>database</code>	string
<code>app.kubernetes.io/part-of</code>	The name of a higher level application this one is part of	<code>wordpress</code>	string
<code>app.kubernetes.io/managed-by</code>	The tool being used to manage the operation of an application	<code>helm</code>	string

Label Examples

Metadata

{Deployment}

Name	Namespace	Creation time	Age
graviteeio-apim3x-portal	gravitee	Nov 11, 2020	3 months

UID
1f6345f5-0013-4026-aa58-44d4c51048d8

Labels

app.kubernetes.io/component: portal	app.kubernetes.io/instance: graviteeio-apim3x	
app.kubernetes.io/managed-by: Helm	app.kubernetes.io/name: apim3	app.kubernetes.io/version: 3.3.3
helm.sh/chart: apim3-3.0.13	Show less	

Annotations

deployment.kubernetes.io/revision: 1	meta.helm.sh/release-name: graviteeio-apim3x
meta.helm.sh/release-namespace: gravitee	

Metadata

{Pod}

Name	Namespace	Creation time	Age
graviteeio-apim3x-portal-869fd8b9-m5kln	gravitee	Nov 11, 2020	3 months

UID
35bc760c-1079-44f9-935f-6734d44e0d74

Labels

app.kubernetes.io/component: portal	app.kubernetes.io/instance: graviteeio-apim3x	app.kubernetes.io/name: apim3
pod-template-hash: 869fd8b9		

Annotations

chaos.alpha.kubernetes.io/enabled: false	checksum/config	cni.projectcalico.org/podIP: 172.20.58.4/32
cni.projectcalico.org/podIPs: 172.20.58.4/32		

Metadata

{Service}

Name	Namespace	Creation time	Age
graviteeio-apim3x-portal	gravitee	Nov 11, 2020	3 months

UID
4672e013-2ca3-437c-a81c-5124e88c56a9

Labels

app.kubernetes.io/component: portal	app.kubernetes.io/instance: graviteeio-apim3x	
app.kubernetes.io/managed-by: Helm	app.kubernetes.io/name: apim3	app.kubernetes.io/version: 3.3.3
helm.sh/chart: apim3-3.0.13	Show less	

Annotations

meta.helm.sh/release-name: graviteeio-apim3x	meta.helm.sh/release-namespace: gravitee
----------------------------------------------	------------------------------------------

Label Command Examples

- `kubectl label -h`
- # Update pod 'foo' with the label 'unhealthy' and the value 'true'.
 - `kubectl label pods foo unhealthy=true`
- # Update all pods in the namespace
 - `kubectl label pods --all status=unhealthy`
- # Update pod 'foo' with the label 'status' and the value 'unhealthy', overwriting any existing value.
 - `kubectl label --overwrite pods foo status=unhealthy`
- # Update pod 'foo' by removing a label named 'bar' if it exists (Does not require the --overwrite flag.)
 - `kubectl label pods foo bar-`
- Some options:
 - `--overwrite=false`
 - `--record=false`

Field Selectors (Selection Via Fields)

- Feature that let us select kubernetes resources based on the value of one or more **resource fields**.
- Supported field selectors vary by Kubernetes resource type.
- All resource types support the **metadata.name** and **metadata.namespace** fields.
- Can use the =, ==, and != operators.
- There are a limited number of fields that can be used for selection.
- Examples of field selector queries:
 - `metadata.name=my-service`
 - `metadata.namespace!=default`
 - `status.phase=Pending`
- Using **--field-selector** option in kubectl command.
 - `kubectl get pods --field-selector status.phase=Running`
 - `kubectl get pods --field-selector=status.phase!=Running,spec.restartPolicy=Always #Chained selectors`
 - `kubectl get statefulsets,services --all-namespaces --field-selector metadata.namespace!=default #Multiple resource types`

Label Selector (Selection Via Labels)

- Labels do not provide uniqueness.
- The label selector is the core grouping primitive in Kubernetes.
- Some Kubernetes objects, such as **replicaset**, also use label selectors to specify sets of pods.

```
apiVersion: apps/v1
kind: Deployment
...
spec:
  replicas: 1
  selector:
    matchLabels:
      app: web-server
```

- We can identify a set of objects by using **--selector** (-l for short) option in kubectl.
 - `kubectl get deploy --selector app=nginx` #Equality-based
 - `kubectl get all --selector app=nginx --show-labels` #Equality-based
 - `kubectl get pods -l 'environment in (production),tier in (frontend)'` #Set-base selector

Two Types of Selectors

- The API currently supports two types of selectors: equality-based and set-based.
- **Equality-based**
 - Equality- or inequality-based requirements allow **filtering by label keys and values**.
 - Three kinds of operators are admitted **=,==,!=**.
 - `environment = production` # selects all resources with key equal to environment and value equal to production
 - `tier != frontend` # selects all resources with key equal to tier and value distinct from frontend, and all resources with no labels with the tier key.
 - `environment=production,tier!=frontend` # using comma
- **Set-based**
 - Allow filtering keys according to a **set of values**.
 - Three kinds of operators are supported: `in`, `notin` and `exists` (only the key identifier).
 - Examples:
 - `environment in (production, qa)` #selects all resources with key equal to environment and value equal to production or qa.
 - `tier notin (frontend, backend)` #selects all resources with key equal to tier and values other than frontend and backend, and all resources with no labels with the tier key.
 - `partition` #selects all resources including a label with key partition; no values are checked.
 - `!partition` #selects all resources without a label with key partition; no values are checked.
 - `partition,environment notin (qa)`

Annotation

- Annotations are not used to **identify** and **select** objects.
- Annotations are used to provide additional information about object.
- Can manually set and automatically set by Kubernetes itself.
- More flexible than label about values format.
- Kubernetes reserves all labels and annotations in the **kubernetes.io** namespace for examples:
 - [alpha.kubernetes.io/provided-node-ip: "10.0.0.1"](#)
 - [ingressclass.kubernetes.io/is-default-class: "true"](#)
- Some examples of information that could be recorded in annotations:
 - Build, release, or image information like timestamps, release IDs, git branch and registry address.
 - Phone or pager numbers of persons responsible.
 - Client libraries or tools information.
- Use **kubectl annotate** command to update the annotations on one or more resources

Annotation Examples

- Show annotation of objects
 - `kubectl get pod <POD NAME> -o jsonpath='{.metadata.annotations}'`
 - `kubectl get deployment nginx-deployment -o yaml | less`
- Set annotation to object examples
 - # Update all pods in the namespace
 - `kubectl annotate pods --all description='my frontend running nginx' # be careful`
 - # Update pod 'foo' with the annotation 'description' and the value 'my frontend running nginx', overwriting any existing value.
 - `kubectl annotate --overwrite pods foo description='my frontend running nginx'`
 - # Update pod 'foo' by removing an annotation named 'description' if it exists.
 - `kubectl annotate pods foo description-`

Metadata as Data and Instructions

