

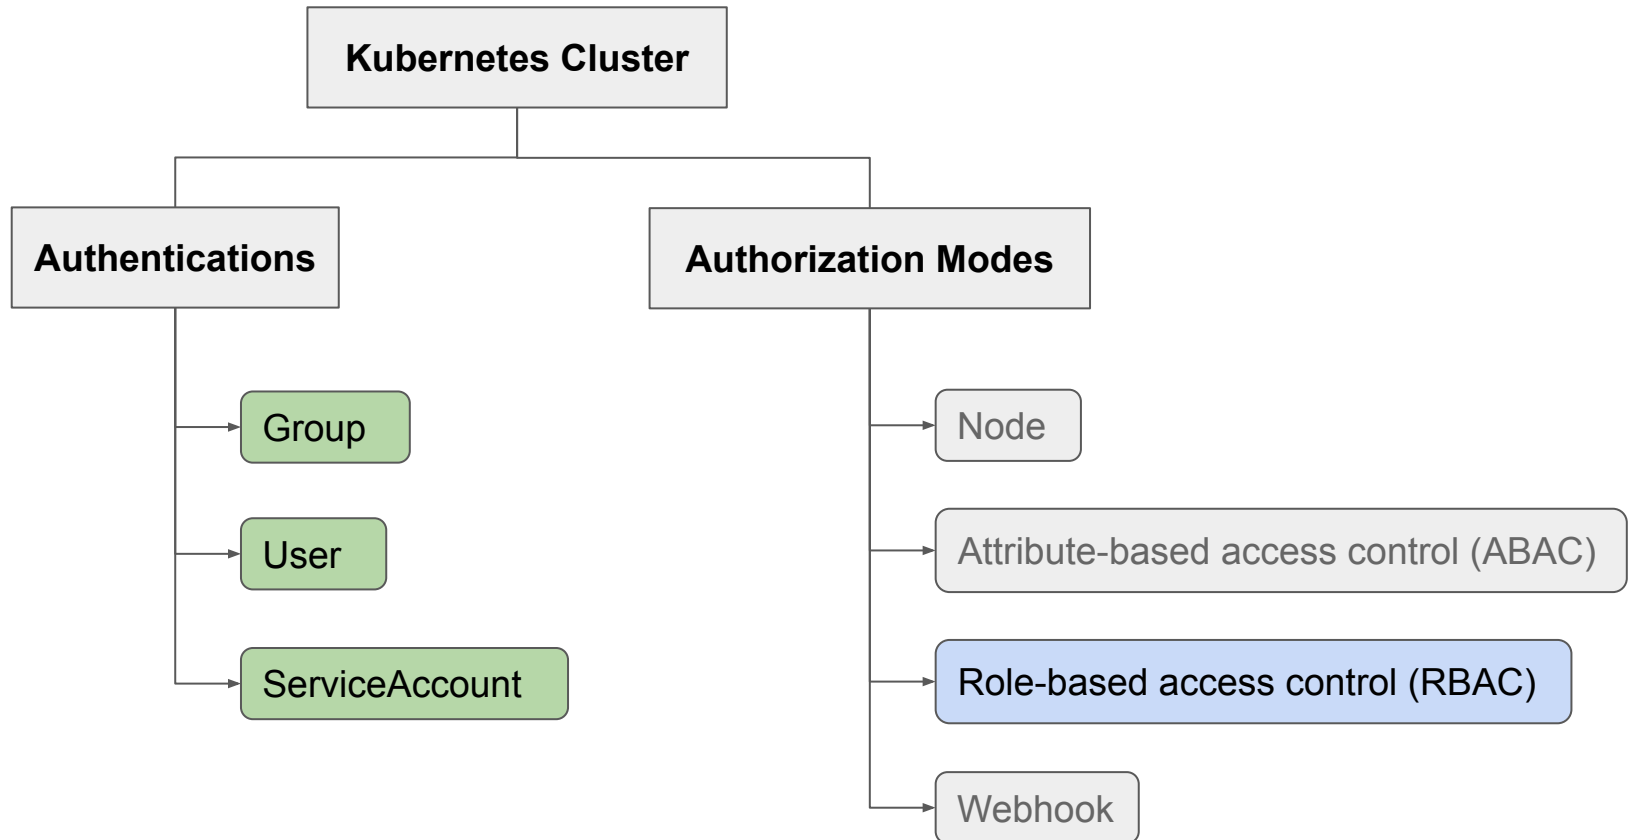
Kubernetes - Part 4

Solar Team

Agenda

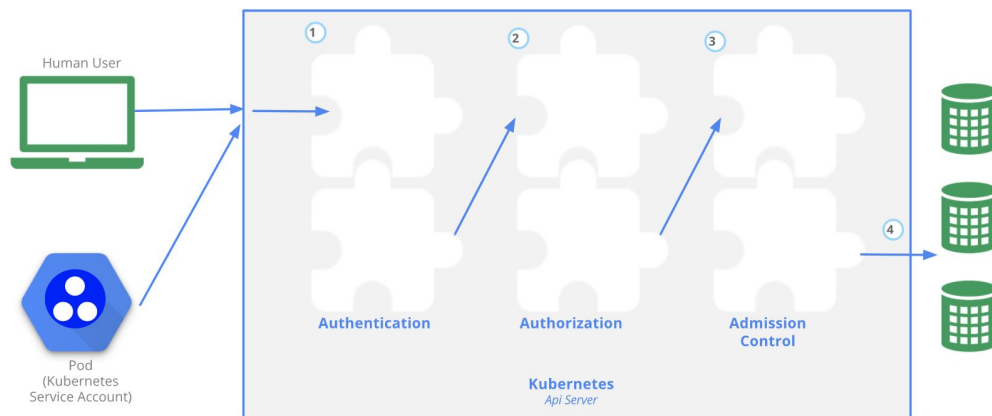
- Kubernetes Authentication
 - User and Group
 - Service Account
- Kubernetes Authorization
 - RBAC Authorization
 - Role
 - Role Binding
 - ClusterRole
 - ClusterRoleBinding

Kubernetes API Access Control



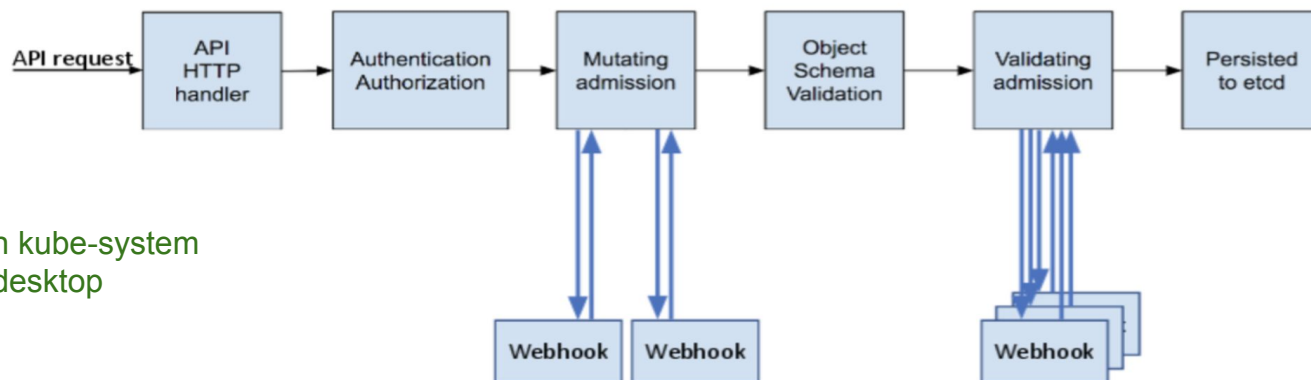
Kubernetes Authentication

- **All actions in a Kubernetes cluster need to be authenticated and authorized.**
- Kubernetes clusters have two categories of users:
 - **User accounts** - for humans (Kubernetes does not have objects which represent normal user accounts.)
 - **Service accounts** - for processes (which run in pods), managed by Kubernetes.
- Support many authentication strategies likes, uses client certificates, bearer tokens, an authenticating proxy, or HTTP basic auth.
- Support anonymous users.



Admission Controllers

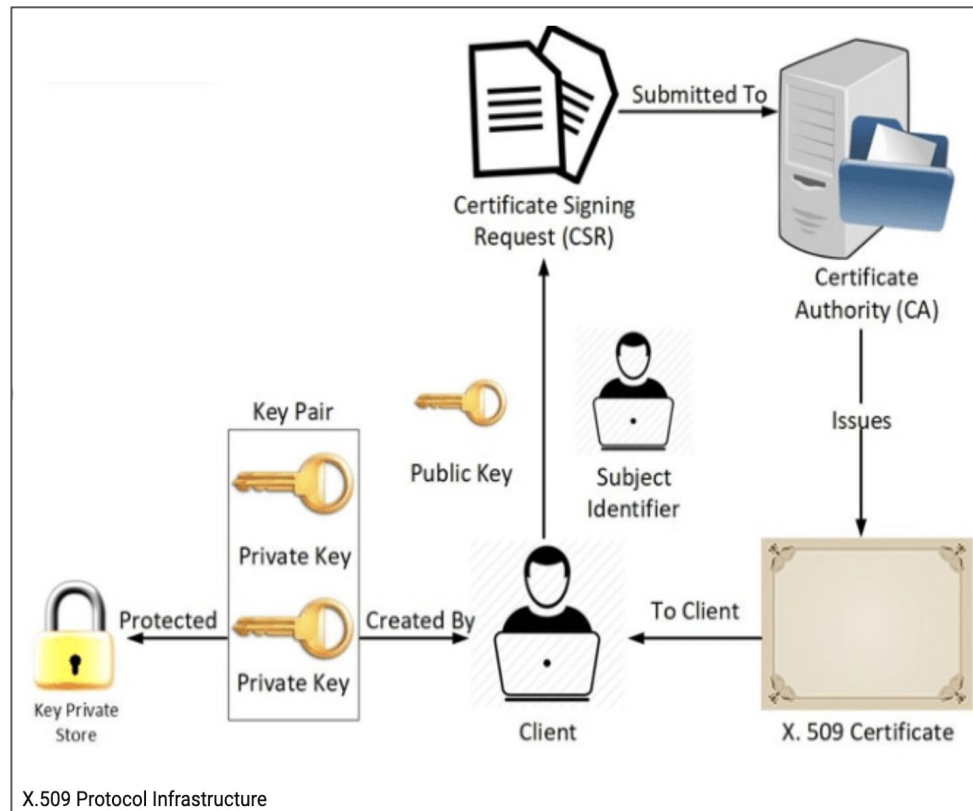
- An admission controller is a piece of code that intercepts requests to the Kubernetes API server prior to persistence of the object, but **after the request is authenticated and authorized**.
- Kubernetes admission controllers are plugins that **govern** and **enforce** how the cluster is used.
- More than 30 admission controllers shipped with Kubernetes.
- Example: **ServiceAccount**, **AlwaysPullImages**, **MutatingAdmissionWebhook** and **ValidatingAdmissionWebhook**.



Admission Controller Phases

- `kubectl describe pods -n kube-system`
`kube-apiserver-docker-desktop`

Create New Kubernetes User - X.509 Client Certificates



Source: https://www.researchgate.net/figure/X509-Protocol-Infrastructure_fig3_320093430

- The Kubernetes Certificates API provide a mechanism to obtain **x509 certificates** by submitting a certificate signing request, and having it asynchronously approved and issued.
- This API can be used to request client certificates to authenticate to kube-apiserver (with the "kubernetes.io/kube-apiserver-client" signerName)
- FEATURE STATE: **Kubernetes v1.19 [stable]**

Creat New Kubernetes User - Prepare CSR

- To create users with [X.509](#) client certificates and how to manage authorizations with the basic Kubernetes [Role-based access control \(RBAC\)](#) API Objects.
- Create the client key
 - `openssl genrsa -out anurak.key 2048`
- Create a [certificate signing request\(CSR\)](#)
 - `openssl req -new -key anurak.key -out anurak.csr`
 - Or `openssl req -new -key anurak.key -out user1.csr -subj "/CN=anurak/O=dev"`
 - **CN** is the [name of the user](#) and **O** is the [group](#) that this user will belong to.

Creat New Kubernetes User - Submit CSR

```
apiVersion: certificates.k8s.io/v1
kind: CertificateSigningRequest
metadata:
  name: anurak
spec:
  groups:
  - system:authenticated
  signerName: kubernetes.io/kube-apiserver-client
  request: LS0tLS1CRUdJTiBDRVJUSUZJQ0FURSBS...Cg==
  usages:
  - client auth
```

csr-anurak.yml

- Create a CertificateSigningRequest and submit it to Kubernetes cluster
- "request" contains an x509 certificate signing request encoded in a "CERTIFICATE REQUEST" PEM block. When serialized as JSON or YAML, the data is additionally base64-encoded.
 - `cat anurak.csr | base64 | tr -d "\n"`
- "kubernetes.io/kube-apiserver-client": issues client certificates that can be used to authenticate to kube-apiserver.
- `kubectl apply -f csr-anurak.yml`

Creat New Kubernetes User - Approve Or Deny CSR

- Approve the CSR:
 - `kubectl certificate approve anurak`
- Or deny the CSR
 - `kubectl certificate deny anurak`
- Retrieve the certificate from the CSR:
 - `kubectl get csr/myuser -o yaml`
- Export the issued certificate from the CertificateSigningRequest. The certificate value is in Base64-encoded format under status.certificate.
 - `kubectl get csr myuser -o jsonpath='{.status.certificate}' | base64 -d > anurak.crt`

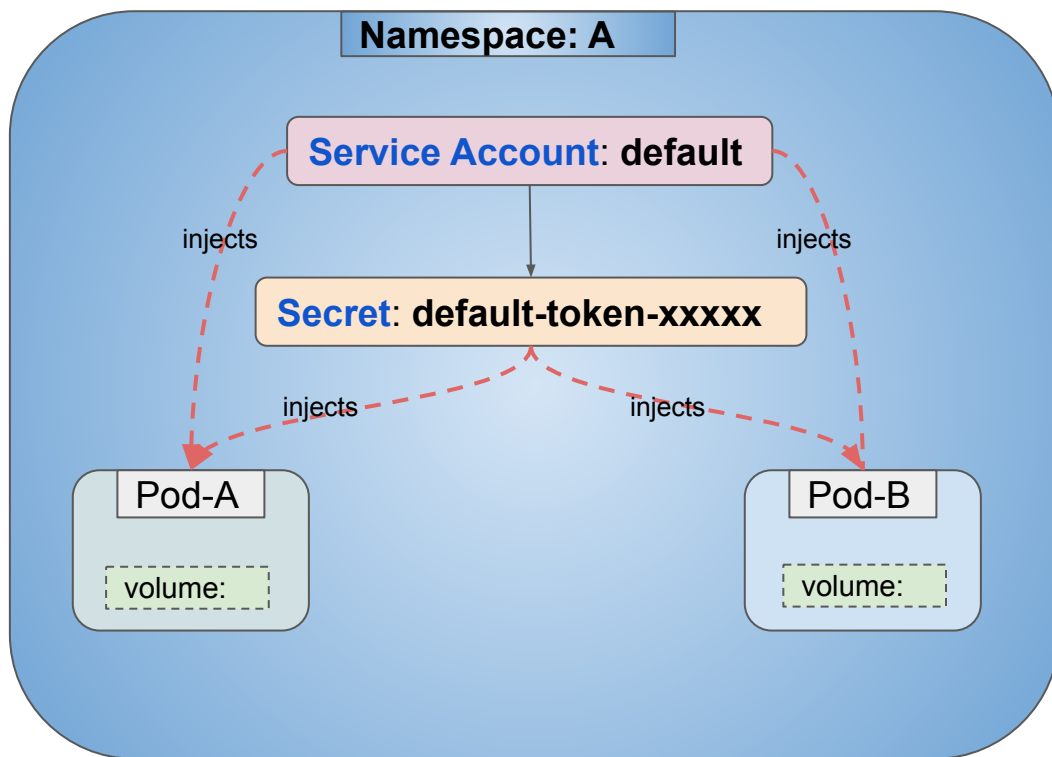
Create New Kubernetes User - Using A New User

- Add this user into the kubeconfig file.
 - `kubectrl config set-credentials anurak --client-key=anurak.key --client-certificate=anurak.crt --embed-certs=true`
- Add the context:
 - `kubectrl config set-context anurak --cluster=kubernetes --user=anurak`
- To test it, change the context to “anurak”:
 - `kubectrl config use-context anurak`
- Create a Role for this new user
 - `kubectrl create role pod-reviewer --verb=get --verb=list --verb=watch --resource=pods`
- Create a RoleBinding for this new user
 - `kubectrl create rolebinding pod-reviewer --role=pod-reviewer --user=anurak`
- Testing
 - `kubectrl auth can-i get pods`
 - `kubectrl get all`
 - `kubectrl get pods`
 - `kubectrl config use-context docker-desktop` # switch user to docker-desktop
 - `kubectrl get po --user=anurak`
 - `kubectrl get rolebinding -owide | grep anurak`

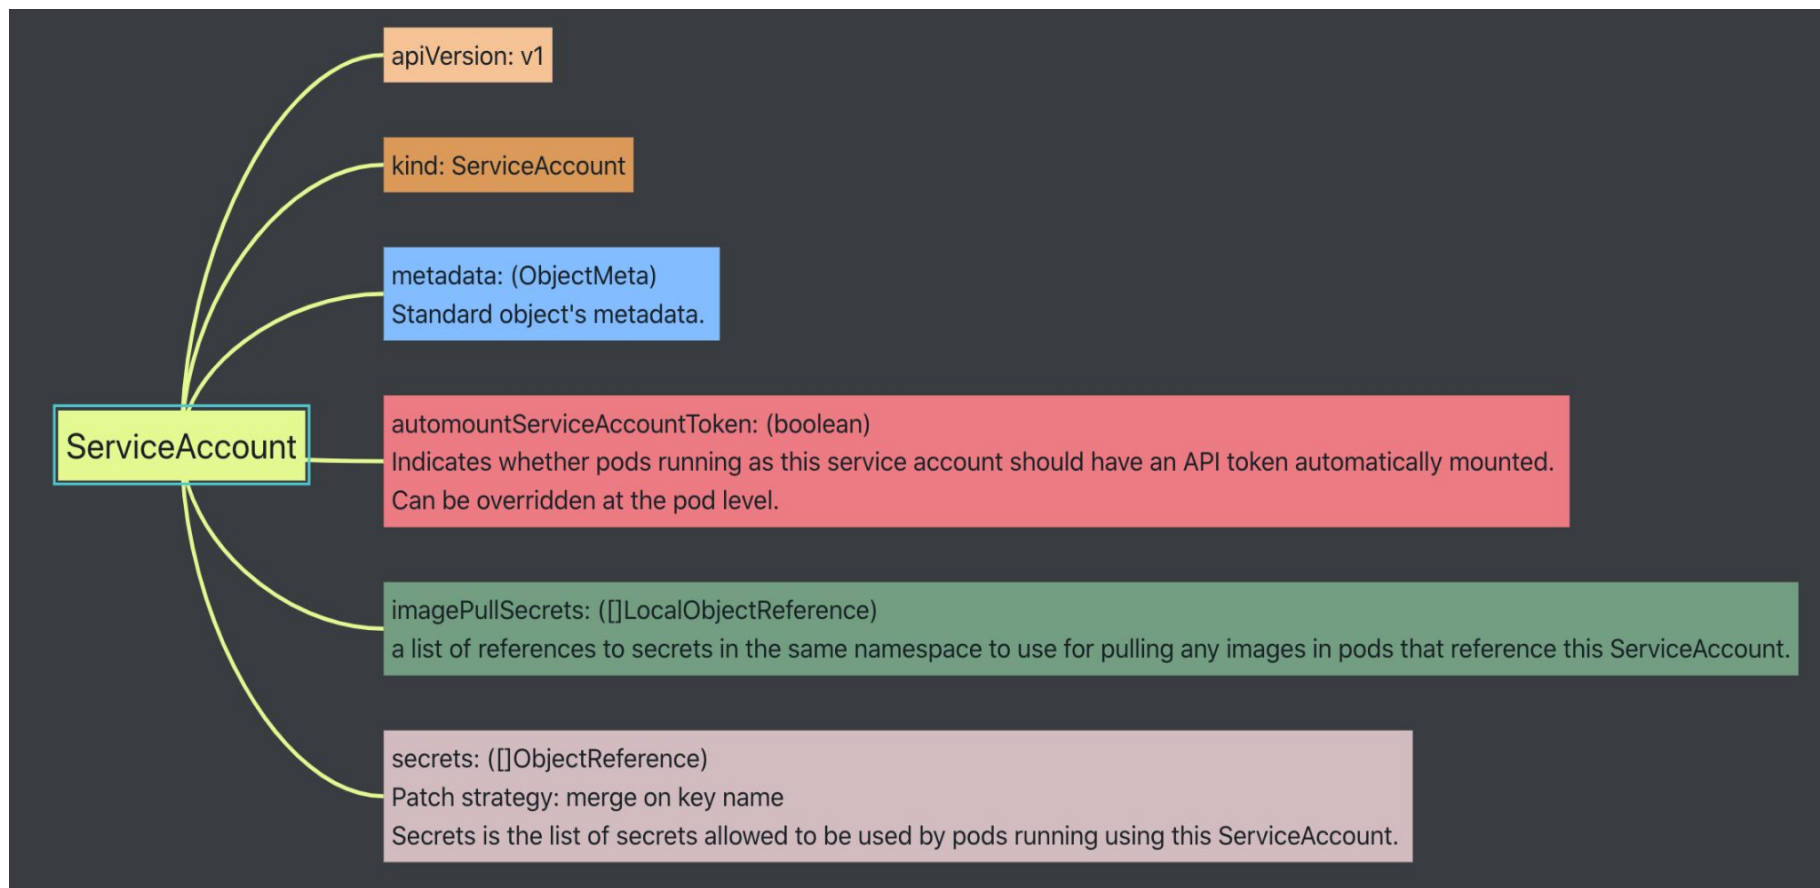
Service Account Overview

- A service account provides an identity for processes that run in a Pod.
- Kubernetes resources that [bound to specific namespaces](#).
- Created automatically by the API server or manually through API calls.
- They are tied to a set of credentials (API credentials) stored as [Secrets](#).
- Every namespace has a default service account resource called [“default”](#).
- A [default service account](#) (if we do not specify a service account) is mounted into pods allowing in-cluster processes to talk to the Kubernetes API.
- Listing service accounts
 - `kubectl get sa`
 - `kubectl get secret`
 - `kubectl get sa default -o yaml`
 - `kubectl describe secret <default-ac-secret>`

Default Service Account



ServiceAccount Object



- `kubectl create secret docker-registry -h`

Default Service Account and Secret

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: default
  namespace: default
  resourceVersion: "330"
secrets:
```

```
- name: default-token-t7ffx
```

```
apiVersion: v1
data:
  ca.crt: LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JS..Cg==
  namespace: ZGVmYXVsdA==
  token:
    ZXlKaGJHY2lPaUpTVXpJMU5pSXNJbXRwWkNklsWklWMGRSYmx..UJB
kind: Secret
metadata:
  annotations:
    kubernetes.io/service-account.name: default
    kubernetes.io/service-account.uid:
    bf597195-7bc3-4f9e-bbba-4e18c58e1ccd
  creationTimestamp: "2021-06-24T01:47:54Z"
  name: default-token-t7ffx
  namespace: default
  resourceVersion: "326"
  type: kubernetes.io/service-account-token
```

Create Service Account

- Create a new service account named my-service-account.
 - `kubectl create serviceaccount anuraksa`
 - `kubectl get sa my-service-account -oyaml`

```
apiVersion: v1
kind: ServiceAccount
metadata:
  creationTimestamp: "2021-05-05T08:46:57Z"
  name: my-service-account
  namespace: default
  resourceVersion: "2613046"
  selfLink: /api/v1/namespaces/default/serviceaccounts/my-service-account
  uid: 72f7d01d-ad7e-11eb-af4b-025000000001
secrets:
- name: my-service-account-token-hs8gz
```

Automounting ServiceAccount Token

- Not every pod needs to contact the API server.
- From version 1.6+ it is possible to **prevent automounting** of service account tokens on pods using `automountServiceAccountToken: false`.
- `automountServiceAccountToken` can be overridden at the Pod level.

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-no-token
  namespace: default
spec:
  containers:
  - image: nginx:1.7.9
    imagePullPolicy: IfNotPresent
    name: nginx
    ports:
    - containerPort: 80
      protocol: TCP
    resources: {}
  serviceAccountName: default
  automountServiceAccountToken: false
```

nginx-pod-no-token.yml

- `kubectl describe po nginx-no-token`
- Checking token file within Pod
 - `kubectl exec -it nginx-no-token --sh`

“cluster-admin” and “docker-for-desktop-binding”

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  annotations:
    rbac.authorization.kubernetes.io/autoupdate: "true"
  creationTimestamp: "2020-12-15T07:59:03Z"
  labels:
    kubernetes.io/bootstrapping: rbac-defaults
name: cluster-admin
resourceVersion: "43"
selfLink: /apis/rbac.authorization.k8s.io/v1/clusterroles/cluster-admin
uid: 65cd40ff-3eab-11eb-95d0-025000000001
rules:
- apiGroups:
  - '*'
  resources:
  - '*'
  verbs:
  - '*'
- nonResourceURLs:
  - '*'
  verbs:
  - '*'
```

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: docker-for-desktop-binding
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: cluster-admin
subjects:
- apiGroup: rbac.authorization.k8s.io
  kind: Group
  name: system:serviceaccounts
  namespace: kube-system
```

Using ServiceAccount Examples

- `kubectl create serviceaccount anuraksa`
- `kubectl run -i --tty --rm=true busybox --image=radial/busyboxplus --serviceaccount=anuraksa --restart=Never -- sh`
 - `TOKEN=$(cat /run/secrets/kubernetes.io/serviceaccount/token)`
 - `curl -H "Authorization: Bearer $TOKEN" https://kubernetes/api/v1/namespaces/default/pods --insecure`
- `kubectl delete clusterrolebinding docker-for-desktop-binding`
- Test curl again.
- `kubectl create role pod-reviewer --verb=get --verb=list --verb=watch --resource=pods`
- `kubectl create rolebinding pod-reviewer --role=pod-reviewer --serviceaccount anuraksa`
- Try again
 - `curl -H "Authorization: Bearer $TOKEN" https://kubernetes/api/v1/namespaces/default/pods --insecure`
- Bind pod-reviewer role binding to user "anurak"
 - `kubectl edit rolebinding pod-reviewer`
 - Add these texts
 - `- apiGroup: rbac.authorization.k8s.io`
 - `kind: User`
 - `name: anurak`

Kubernetes Authorization

- Determine whether a request is [allowed or denied](#).
- Multiple authorization modules can be configured and are checked in sequence.
- Authorization Modes.
 - [Node](#) - A special-purpose authorization mode that grants permissions to kubelets based on the pods they are scheduled to run.
 - [ABAC](#) - Attribute-based access control (ABAC) defines an access control paradigm.
 - [RBAC](#) - Role-based access control (RBAC) is a method of regulating access to computer or network resources based on the roles of individual users within an enterprise.
 - [Webhook](#) - A WebHook is an HTTP callback: an HTTP POST that occurs when something happens.
- Checking current cluster authorization modes:
 - `kubectl cluster-info dump | grep authorization-mode`
 - `kubectl api-versions`
- Checking Kubernetes API access.
 - `kubectl auth can-i create deployments --namespace kube-system`
 - `kubectl auth can-i create deployments --namespace kube-system --as anurak`

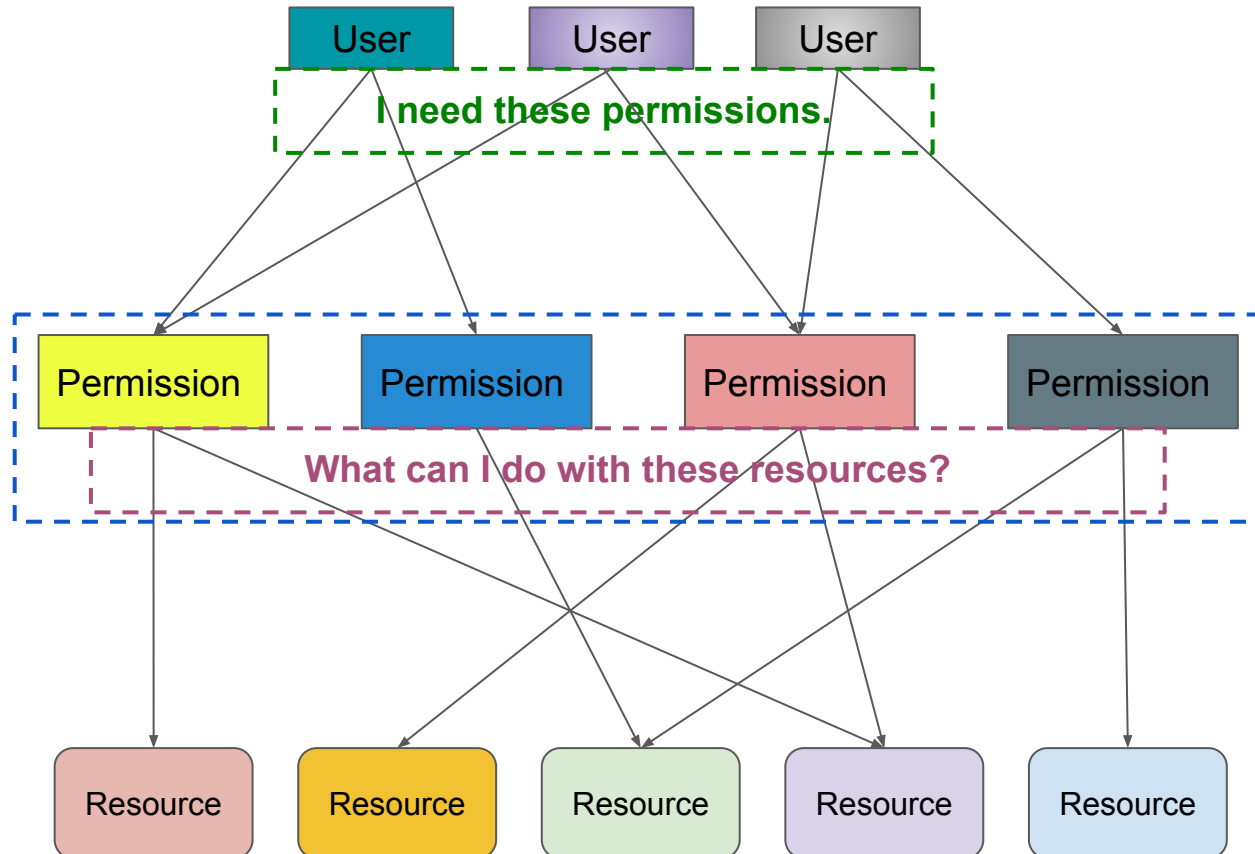
Authorization Modes

- Authorization needs to set up for a new service account.

Authorization Mode	Source	Usage
Node	API Server built-in	kubelet (Internal uses)
ABAC (Attribute-based access control)	local file	Insecure, deprecated
RBAC (Role-based access control)	API Objects	Users and administrators
Webhook	External services	Integration
AlwaysDeny AlwaysAllow	API Server built-in	Testing

- Only RBAC and webhook are production-grade authorization.

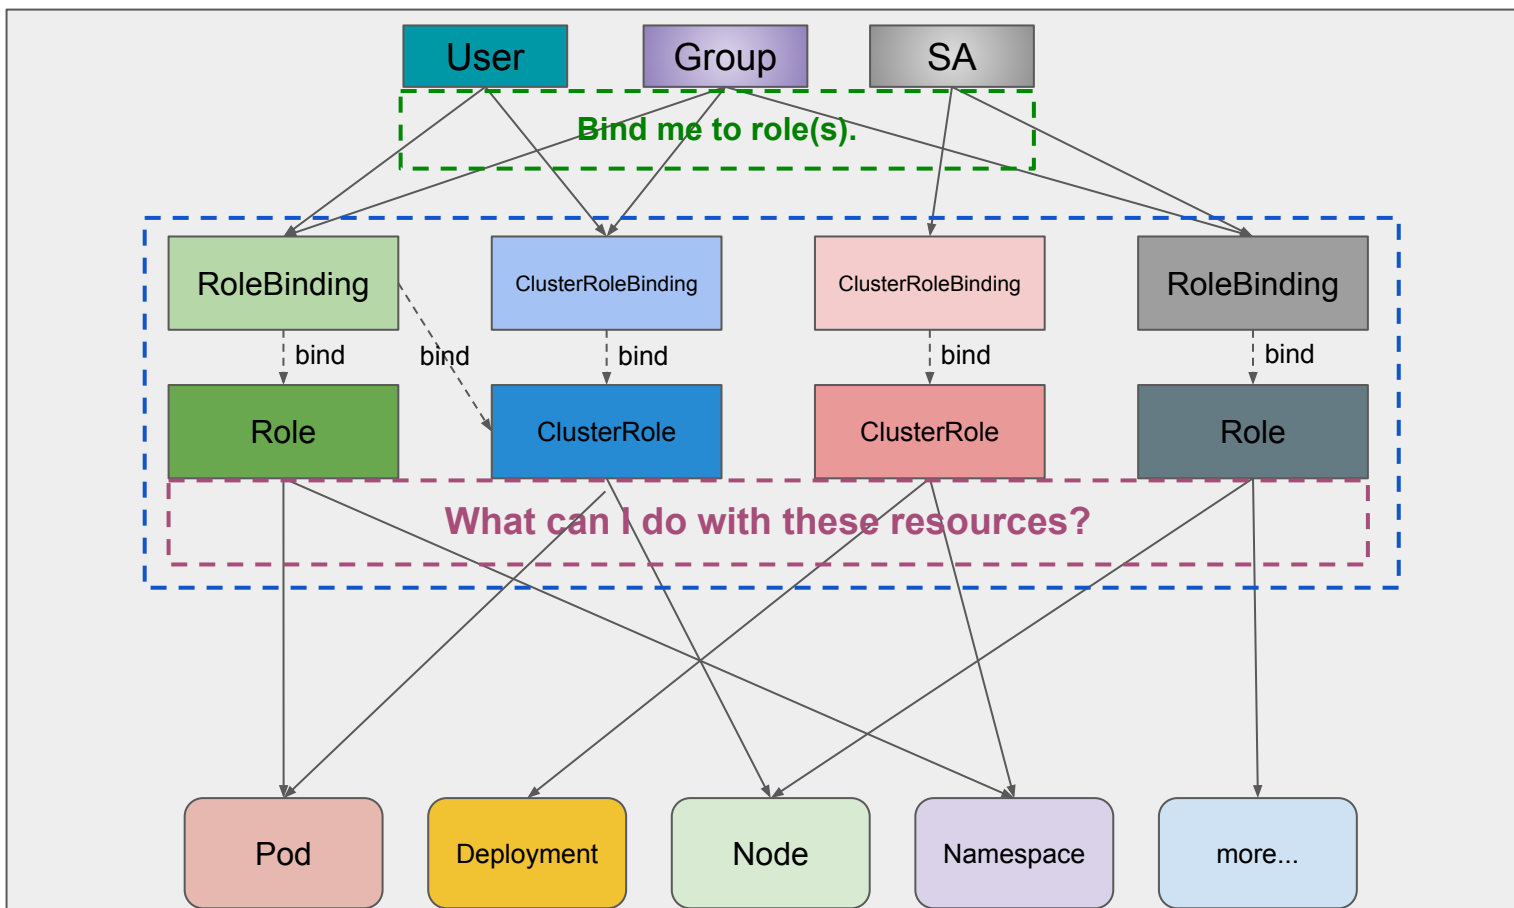
Permission Base Security



RBAC Authorization

- Kubernetes RBAC is a way to define **which users can do what within a Kubernetes cluster**.
- To enable RBAC, start the API server with the `--authorization-mode` flag
 - `kube-apiserver --authorization-mode=RBAC`
 - Check current config: `kubectl cluster-info dump | grep authorization-mode`
- The RBAC API declares **four kinds** of Kubernetes object: **Role**, **ClusterRole**, **RoleBinding** and **ClusterRoleBinding**.
- **Role or ClusterRole** contains rules that represent a **set of permissions**.
- **RoleBinding or ClusterRoleBinding** grants the permissions defined in a role to a user or set of users.
- If you want to **define a role within a namespace**, use a **Role**; if you want to **define a role cluster-wide**, use a **ClusterRole**.

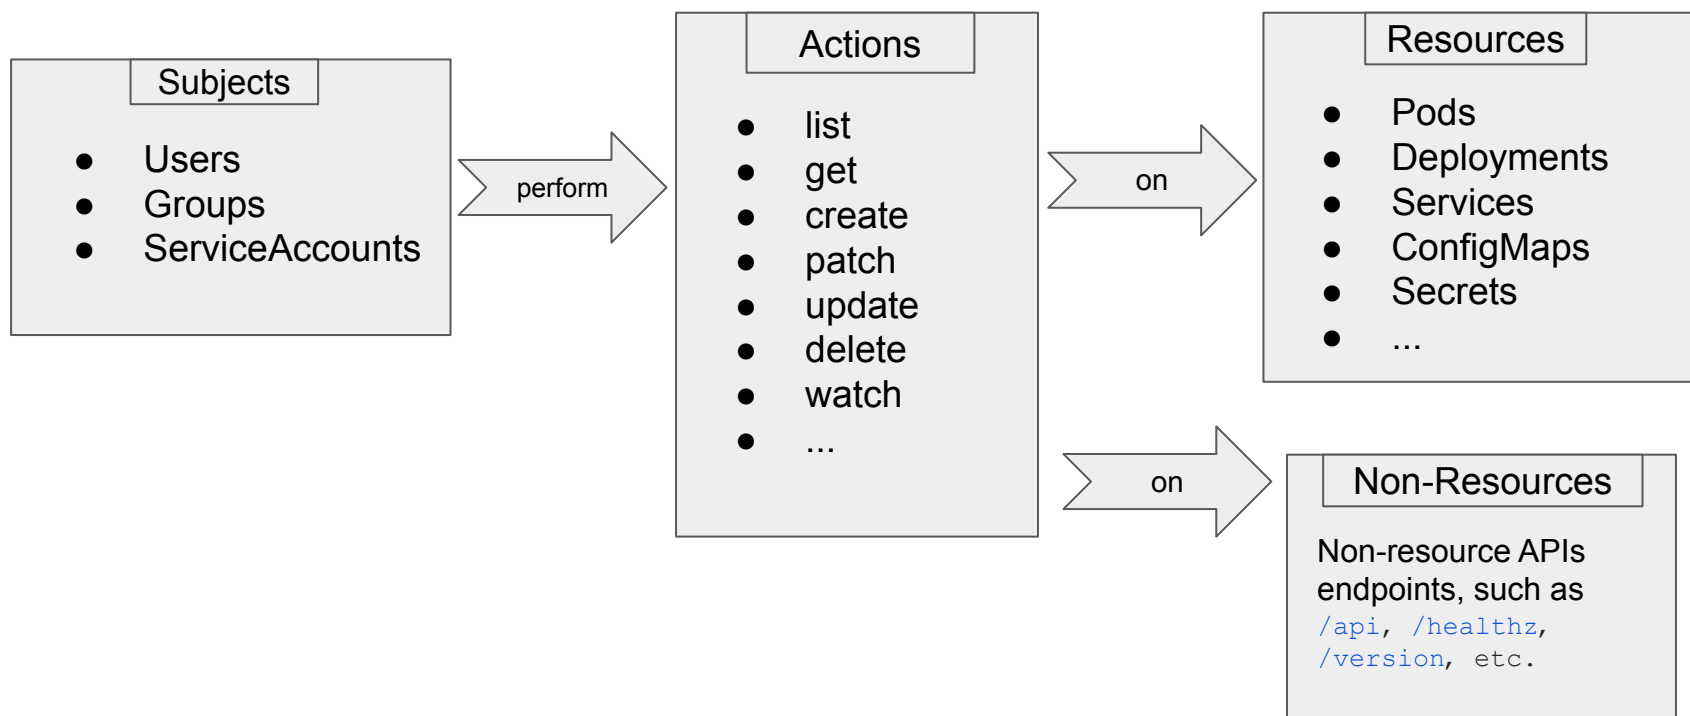
Kubernetes RBAC Authorization



- An RBAC **Role** or **ClusterRole** contains rules that represent a set of permissions.
- A **role binding** grants the permissions defined in a role to a user or set of users.

RBAC Concept - Subjects, Actions and Resources

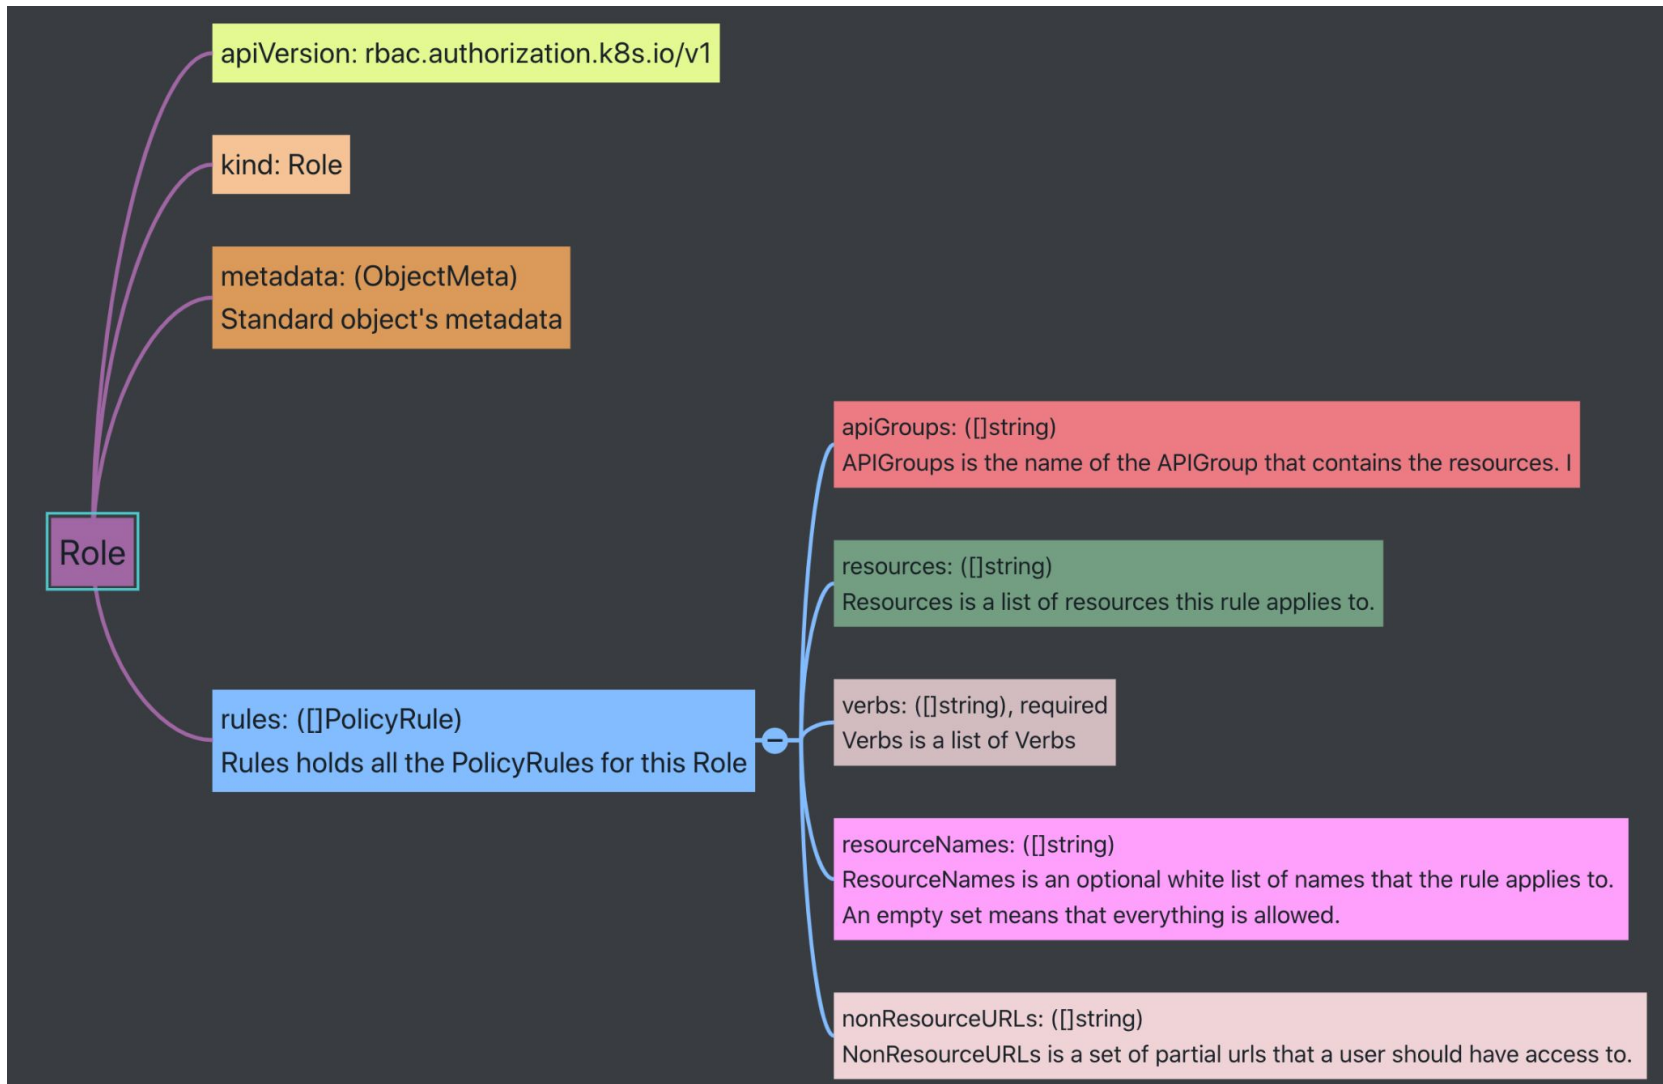
- RBAC allows **subjects** to perform certain **actions** on certain **resource** types.
- Resource actions
 - `kubectl api-resources -owide`



Role

- **Role** - define the **fine-grained permissions** that specify a set of **resources** and **actions** (verbs) that roles can access and manipulate.
- The permissions can only be granted for the resources that are in the **same namespace as the Roles**.
- Permissions are purely additive (there are no "deny" rules)
- **Rules** in Role define a set of permissions for certain operations and resources by specifying **apiGroups, resources, and verbs**.

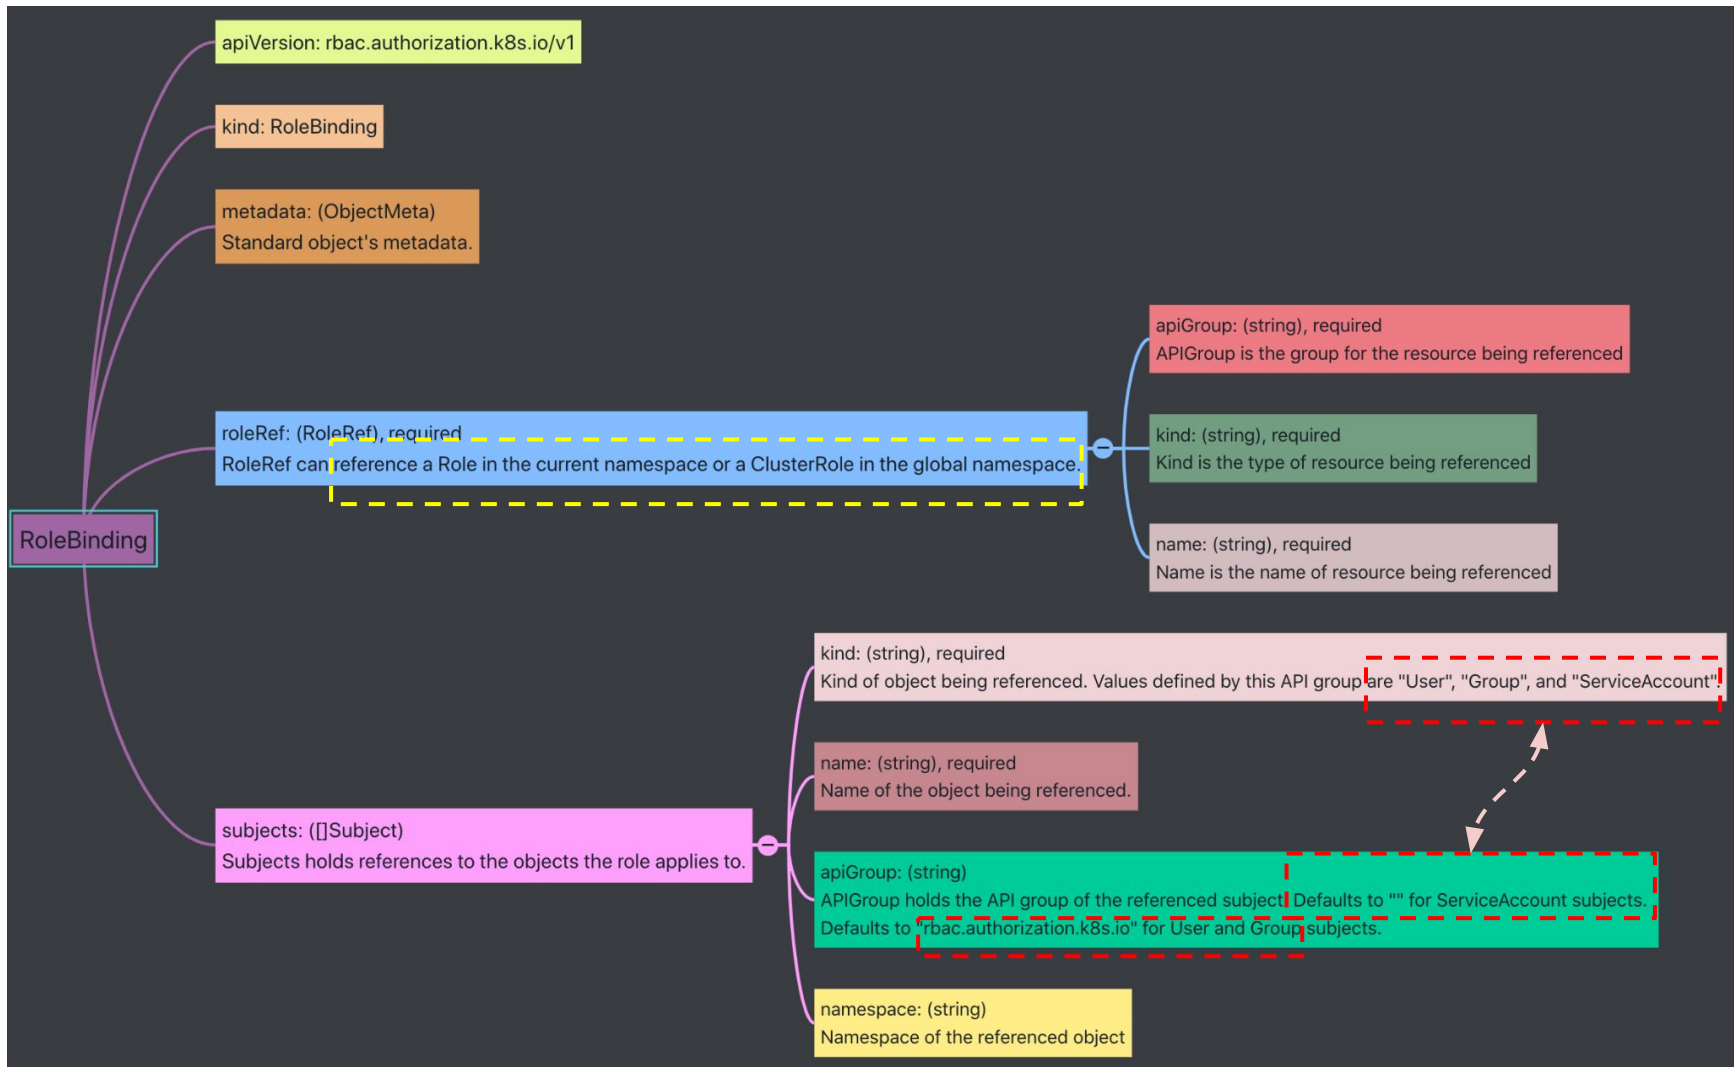
Role Object



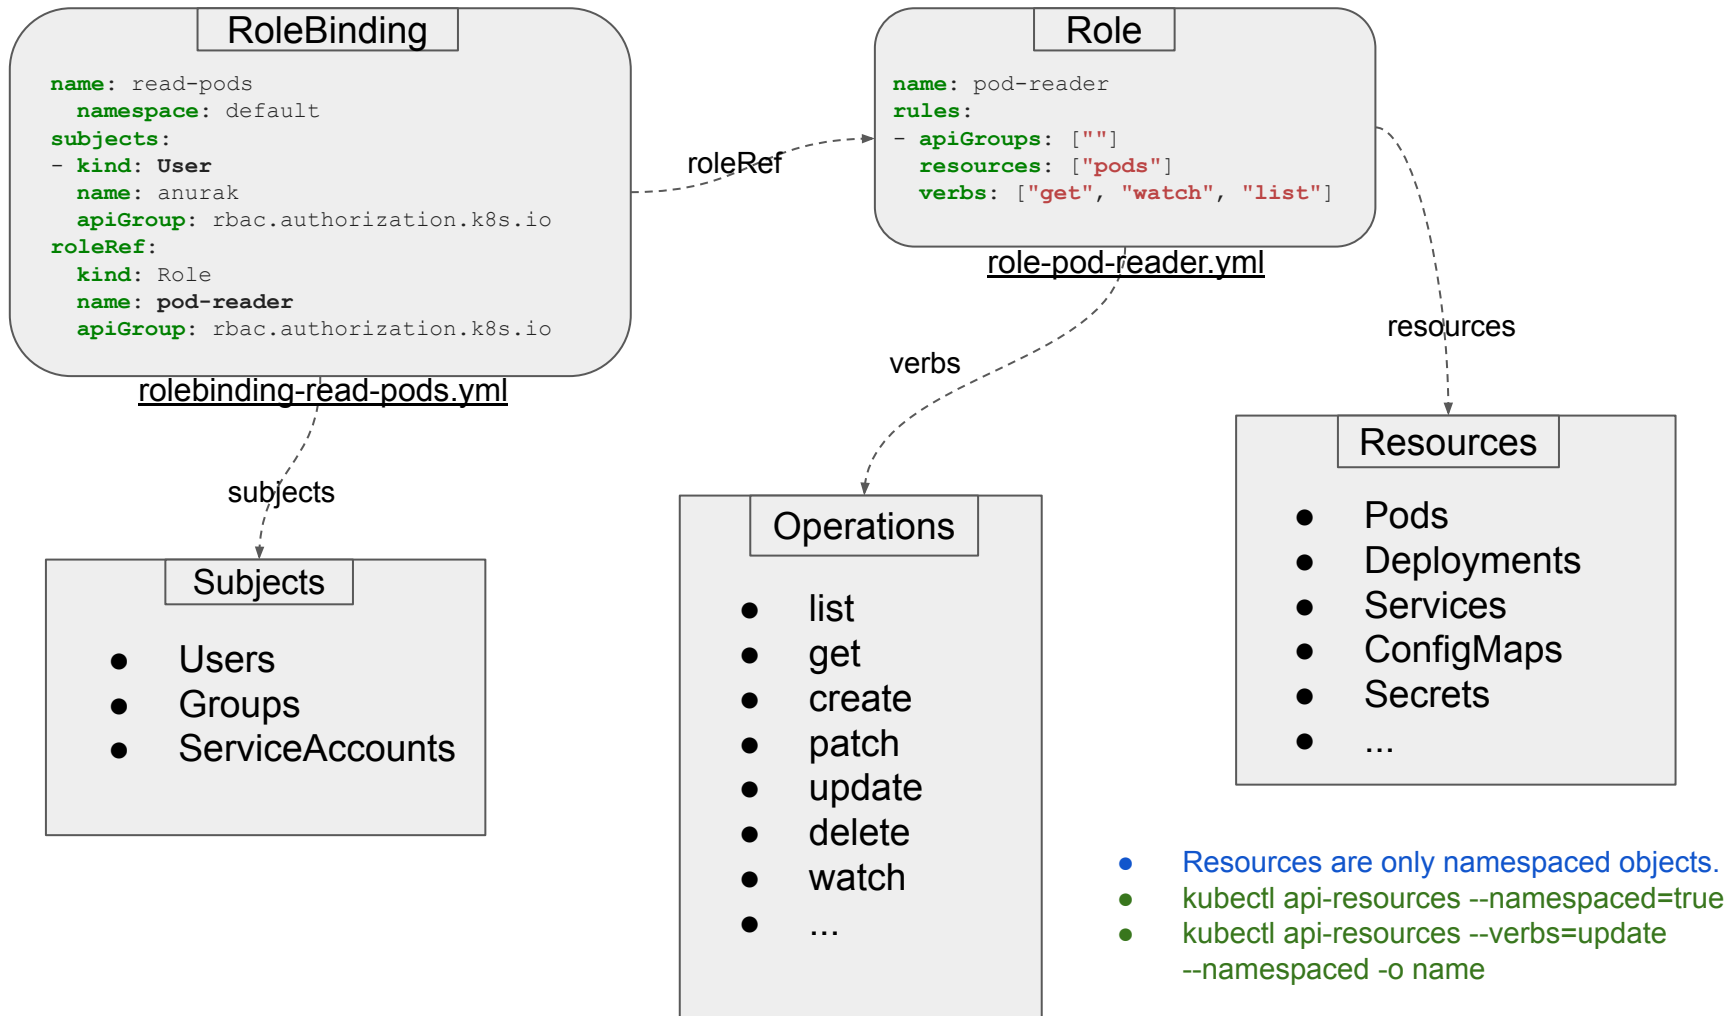
RoleBinding

- **RoleBinding** - grants the permissions defined in a role to a user or set of users.
- A RoleBinding holds a list of subjects (users, groups, or service accounts), and a reference to the role being granted.
- A RoleBinding may reference any Role in the same namespace.
- Alternatively, a RoleBinding can reference a ClusterRole and bind that ClusterRole to the namespace of the RoleBinding

RoleBinding Object



RBAC Roles and RoleBinding



Note: Check the resource and verb with this command.: `kubectl api-resources -owide`

Create Role Using kubectl Examples

- Create a Role named "pod-reader" that allows users to perform get, watch and list on pods:
 - `kubectl create role pod-reader --verb=get --verb=list --verb=watch --resource=pods`
- Create a Role named "pod-reader" with resourceNames specified:
 - `kubectl create role pod-reader --verb=get --resource=pods --resource-name=readablepod --resource-name=anotherpod`
- Create a Role named "foo" with subresource permissions:
 - `kubectl create role foo --verb=get,list,watch --resource=pods,pods/status`
- Create a Role named "foo" with apiGroups specified:
 - `kubectl create role foo --verb=get,list,watch --resource=replicasets.apps`
- Get Help
 - `kubectl create role -h`

Create RoleBinding Using kubectl Examples

- Within the namespace "acme", grant the permissions in the "admin" ClusterRole to a user named "bob":
 - `kubectl create rolebinding bob-admin-binding --clusterrole=admin --user=bob --namespace=acme`
- Within the namespace "acme", grant the permissions in the "view" ClusterRole to the service account in the namespace "acme" named "myapp":
 - `kubectl create rolebinding myapp-view-binding --clusterrole=view --serviceaccount=acme:myapp --namespace=acme`
- Within the namespace "acme", grant the permissions in the "view" ClusterRole to a service account in the namespace "myappnamespace" named "myapp":
 - `kubectl create rolebinding myappnamespace-myapp-view-binding --clusterrole=view --serviceaccount=myappnamespace:myapp --namespace=acme`

Role and RoleBinding YAML Example

```
apiVersion:
rbac.authorization.k8s.io/v1
kind: Role
metadata:
  namespace: default
  name: pod-reader
rules:
- apiGroups: [""] # "" indicates
the core API group
  resources: ["pods"]
  verbs: ["get", "watch", "list"]
```

role-pod-reader.yml

```
- kind: Group
  name: <GROUP NAME>
  apiGroup: rbac.authorization.k8s.io
```

```
- kind: ServiceAccount
  name: <Service Account NAME>
  apiGroup: ""
```

```
apiVersion: rbac.authorization.k8s.io/v1
# This role binding allows "jane" to read pods in the
"default" namespace.
# You need to already have a Role named "pod-reader" in that
namespace.
kind: RoleBinding
metadata:
  name: read-pods
  namespace: default
subjects:
# You can specify more than one "subject"
- kind: User
  name: anurak # "name" is case sensitive
  apiGroup: rbac.authorization.k8s.io
roleRef:
  # "roleRef" specifies the binding to a Role / ClusterRole
  kind: Role #this must be Role or ClusterRole
  name: pod-reader # this must match the name of the Role or
ClusterRole you wish to bind to
  apiGroup: rbac.authorization.k8s.io
```

rolebinding-read-pods.yml

ClusterRole and ClusterRoleBinding

- ClusterRole and ClusterRoleBinding, by contrast, are non-namespaced resources.
- They are used to grant permissions to cluster-wide resources
- Use ClusterRole to grant access to:
 - cluster-scoped resources (like nodes)
 - non-resource endpoints (like /healthz)
 - namespaced resources (like Pods), across all namespaces
- You can **aggregate** several ClusterRoles into one combined ClusterRole.

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: monitoring
aggregationRule:
  clusterRoleSelectors:
    - matchLabels:
        rbac.example.com/aggregate-to-monitoring: "true"
rules: [] # The control plane automatically fills in the rules
```

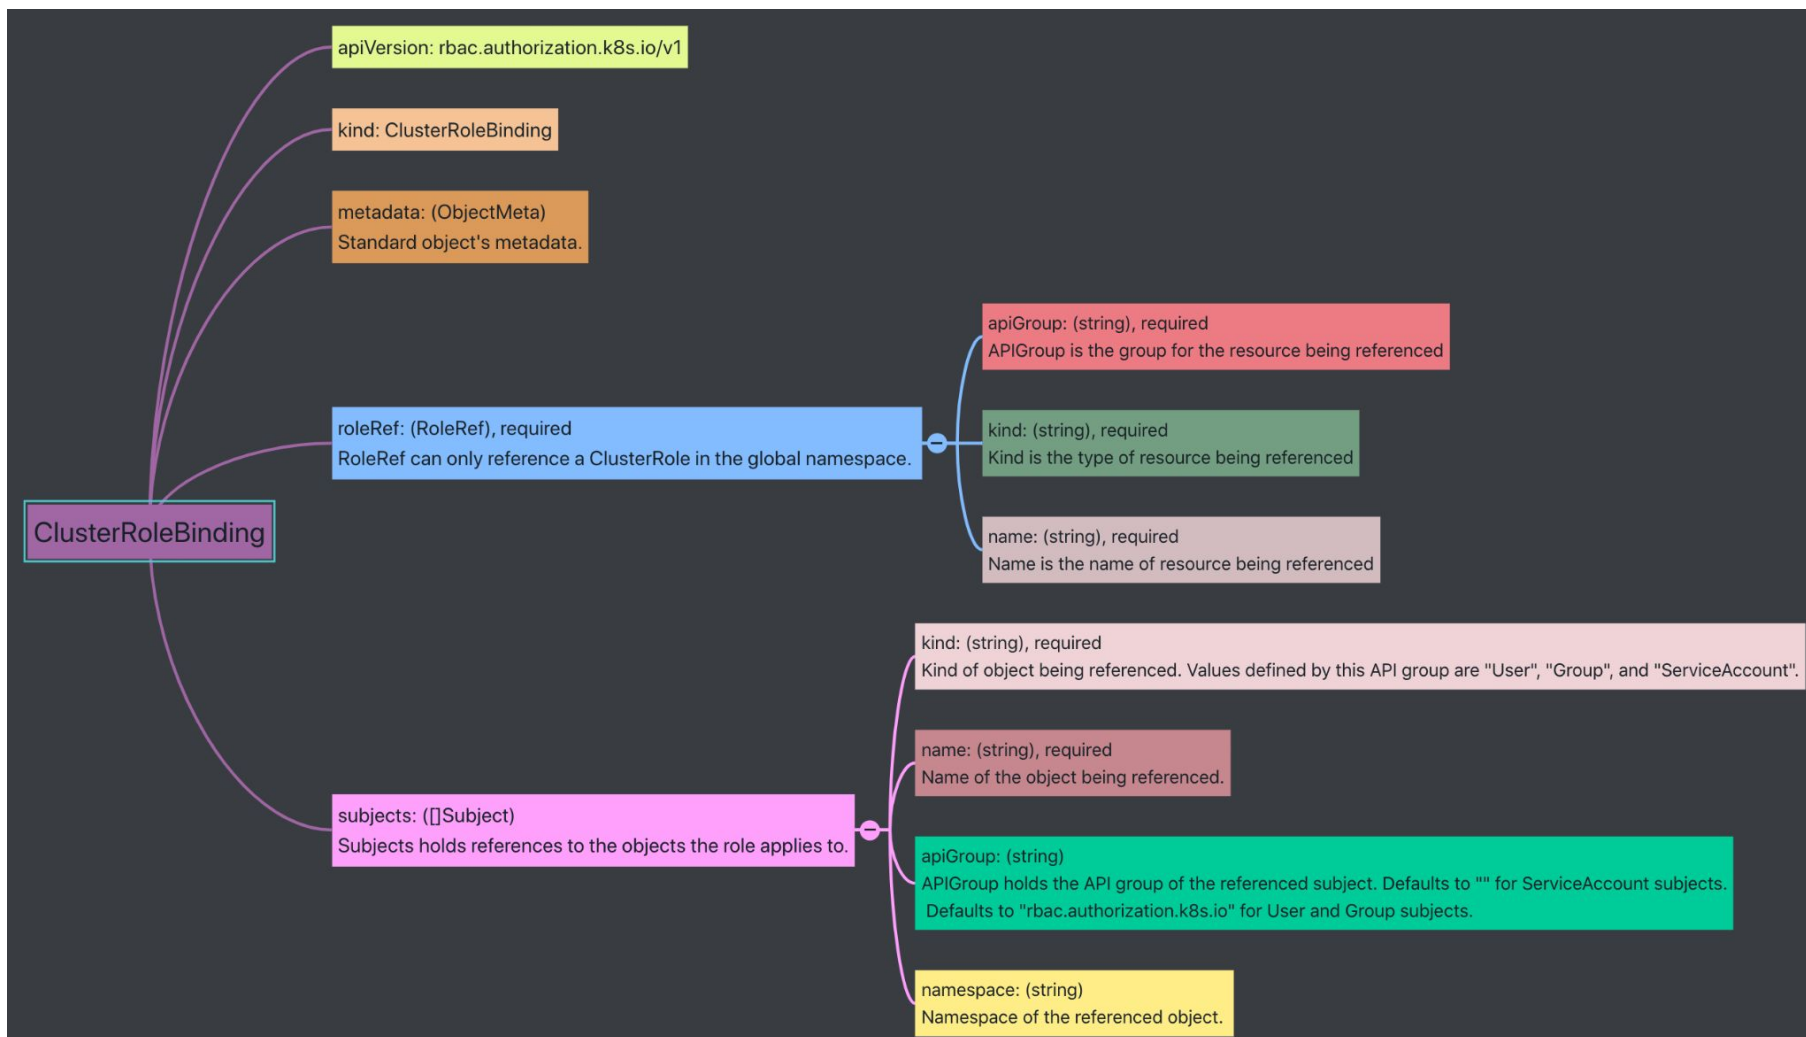
Default ClusterRole and ClusterRoleBinding

- API servers create a set of default ClusterRole and ClusterRoleBinding objects.
- Prefix with “**system:**”
- All of the default ClusterRoles and ClusterRoleBindings are labeled with kubernetes.io/bootstrapping=rbac-defaults.
- `kubectl get clusterrole -l kubernetes.io/bootstrapping=rbac-defaults`
- Some of the default ClusterRoles are not **system:** prefixed.
 - `cluster-admin` - Allows super-user access to perform any action on any resource.
 - `admin` - Allows admin access, intended to be granted within a namespace using a RoleBinding.
 - `edit` - Allows read/write access to most objects in a namespace. This role does not allow viewing or modifying roles or role bindings.
 - `view` - Allows read-only access to see most objects in a namespace. It does not allow viewing roles or role bindings.
 - `kubectl get clusterrole cluster-admin -oyaml`
 - `kubectl get clusterrolebinding cluster-admin -oyaml`

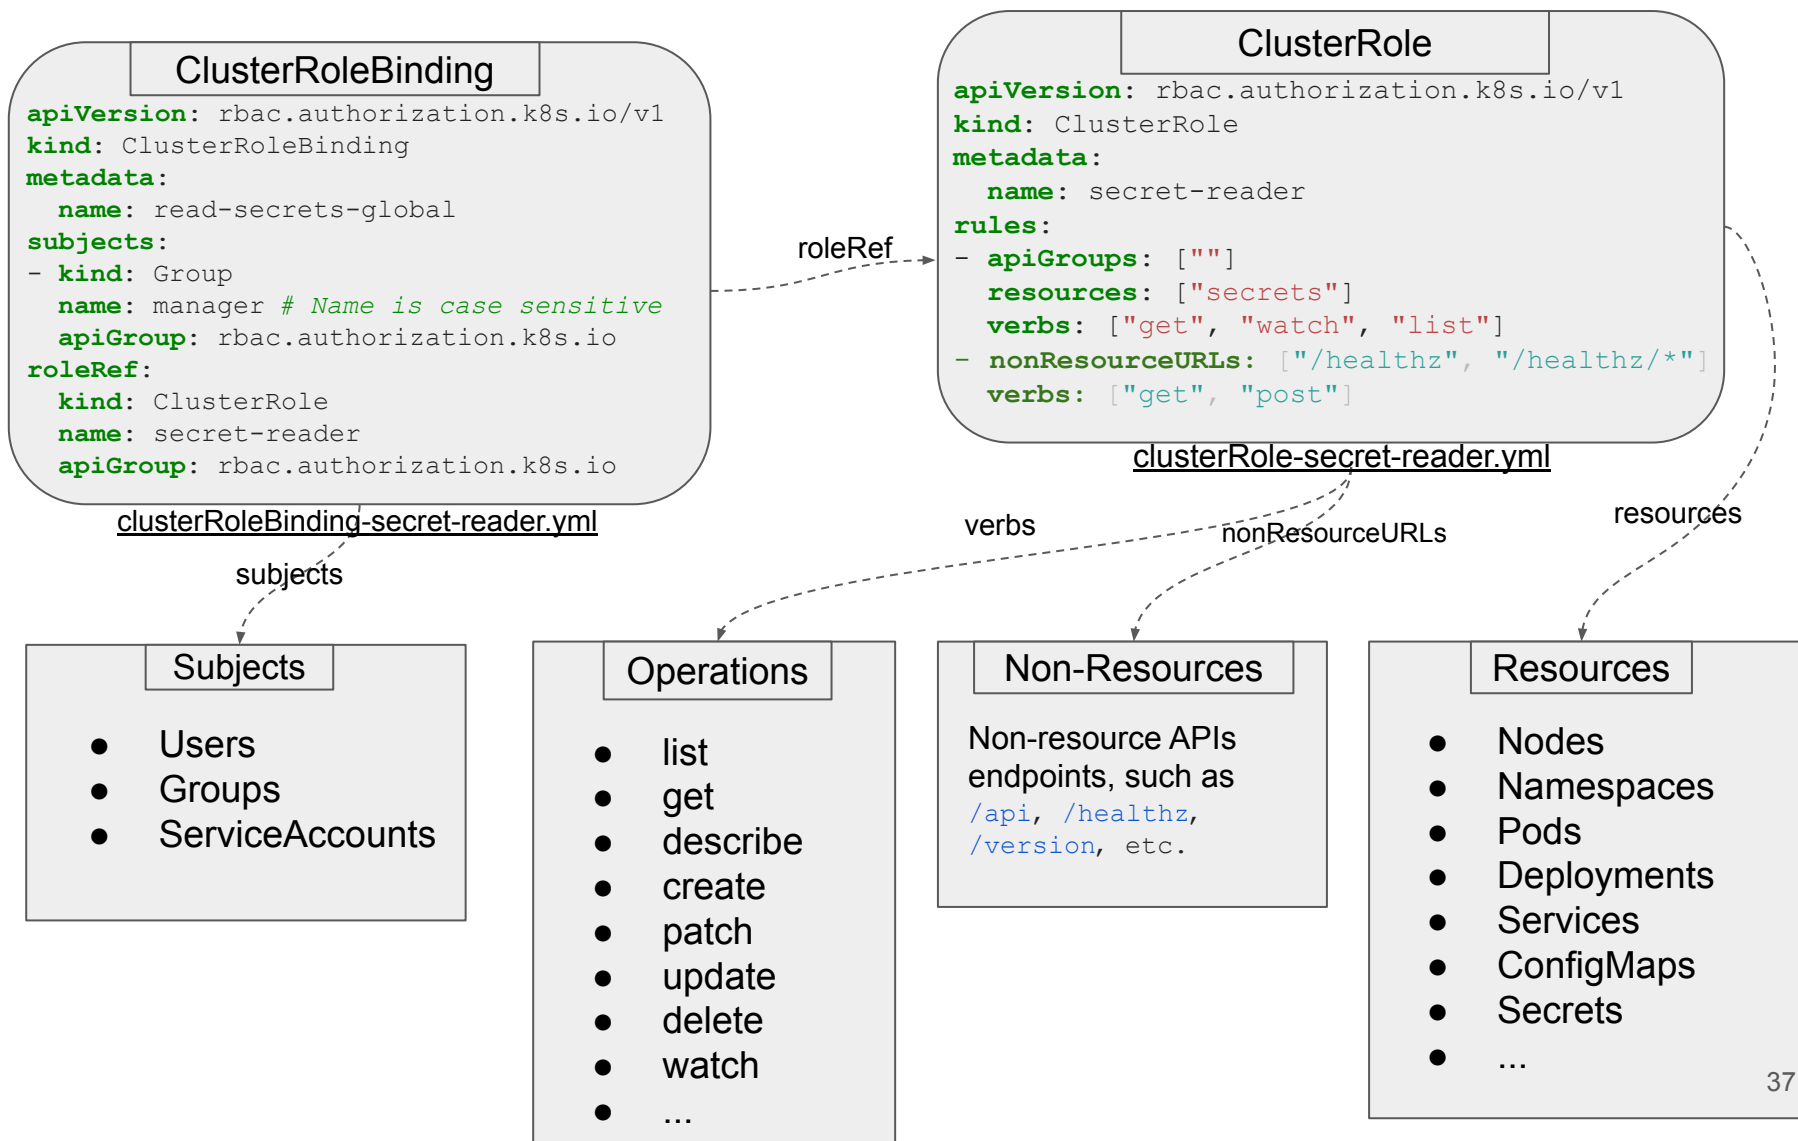
ClusterRole Object



ClusterRoleBinding Object



RBAC ClusterRoles and ClusterRoleBinding



ClusterRole Examples

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  # "namespace" omitted since ClusterRoles are not
  namespace
  name: secret-reader
rules:
- apiGroups: [""]
  resources: ["secrets"]
  verbs: ["get", "watch", "list"]
- nonResourceURLs: ["/healthz", "/healthz/*"]
  verbs: ["get", "post"]
```

clusterRole-secret-reader.yml

- `kubectl apply -f clusterRole-secret-reader.yml`
- `kubectl create rolebinding secret-reader --clusterrole=secret-reader --serviceaccount default:anuraksa --namespace=default`
- Within busybox shell
 - `curl -H "Authorization: Bearer $TOKEN" https://kubernetes/healthz?verbose --insecure`

Create ClusterRole Using kubectl Exaples

- Create a ClusterRole named "pod-reader" that allows user to perform get, watch and list on pods:
 - `kubectl create clusterrole pod-reader --verb=get,list,watch --resource=pods`
- Create a ClusterRole named "pod-reader" with resourceNames specified:
 - `kubectl create clusterrole pod-reader --verb=get --resource=pods --resource-name=readablepod --resource-name=anotherpod`
- Create a ClusterRole named "foo" with apiGroups specified:
 - `kubectl create clusterrole foo --verb=get,list,watch --resource=replicasets.apps`
- Create a ClusterRole named "foo" with nonResourceURL specified:
 - `kubectl create clusterrole "foo" --verb=get --non-resource-url=/logs/*`
- Create a ClusterRole named "monitoring" with an aggregationRule specified:
 - `kubectl create clusterrole monitoring --aggregation-rule="rbac.example.com/aggregate-to-monitoring=true"`

Create ClusterRoleBinding Using kubectl Examples

- Across the entire cluster, grant the permissions in the "cluster-admin" ClusterRole to a user named "root":
 - `kubectl create clusterrolebinding root-cluster-admin-binding --clusterrole=cluster-admin --user=root`
- Across the entire cluster, grant the permissions in the "system:node-proxier" ClusterRole to a user named "system:kube-proxy":
 - `kubectl create clusterrolebinding kube-proxy-binding --clusterrole=system:node-proxier --user=system:kube-proxy`
- Across the entire cluster, grant the permissions in the "view" ClusterRole to a service account named "myapp" in the namespace "acme":
 - `kubectl create clusterrolebinding myapp-view-binding --clusterrole=view --serviceaccount=acme:myapp`

End.