

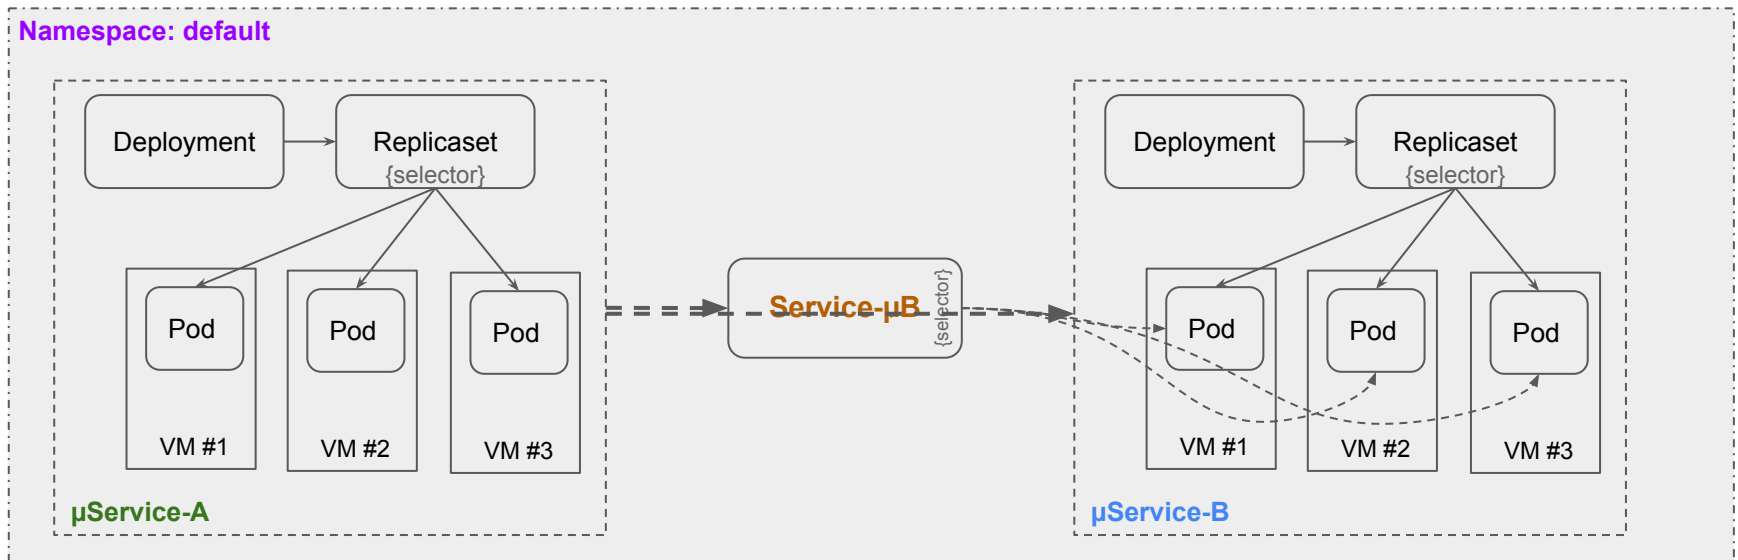
Kubernetes - Part 5

Solar Team

Agenda

- Kubernetes Services
 - ClusterIP
 - NodePort
 - LoadBalancer
 - ExternalName
- Ingress

Service is an Abstraction!



- A Pod is not intended to be treated as a durable, long-lived entity.
- Pods have their own IP address.
- The Service abstraction enables the **decoupling** communication between microservices.

Services Overview

- Services are an **abstract way to expose an application** running on a set of Pods **as a network service**.
- Kubernetes uses a **single DNS name** for a set of pods and can **load-balance** across them.
- Service uses a **label selector** to determine pods.
- Services provide discoverable names and load balancing to Pod replicas.

Exposing Service Basic Example

- Using `--expose` option when creating a Pod.
 - `kubectl run nginx --image=nginx --restart=Never --port=80 --expose`
- Uses `kubectl expose` command to expose service from existing Pod(s).
 - Possible resources include `pod`, `service`, `deployment`, `replicaset`.
 - Create deployment
 - `kubectl run nginx --image=nginx`
 - Or, `kubectl create deploy nginx --image=nginx`
 - Expose deployment as a service
 - `kubectl expose deploy nginx --port=8080 --target-port=80`
 - `--port=` `port that the service should serve on`.
 - `--target-port=` `port on the container` that the service should direct traffic to.
- Testing service with `cURL`

Exposing Service - Service Selector

```
✓ kubectl describe po nginx
Name:          nginx-65f88748fd-rt2h4
Namespace:     default
Priority:       0
Node:          docker-desktop/192.168.65.3
Start Time:    Mon, 17 May 2021 09:17:17 +0700
Labels:        app=nginx
               pod-template-hash=65f88748fd
Annotations:   <none>
Status:        Running
IP:            10.1.2.26
IPs:           <none>
Controlled By: ReplicaSet/nginx-65f88748fd
Containers:
  nginx:
```

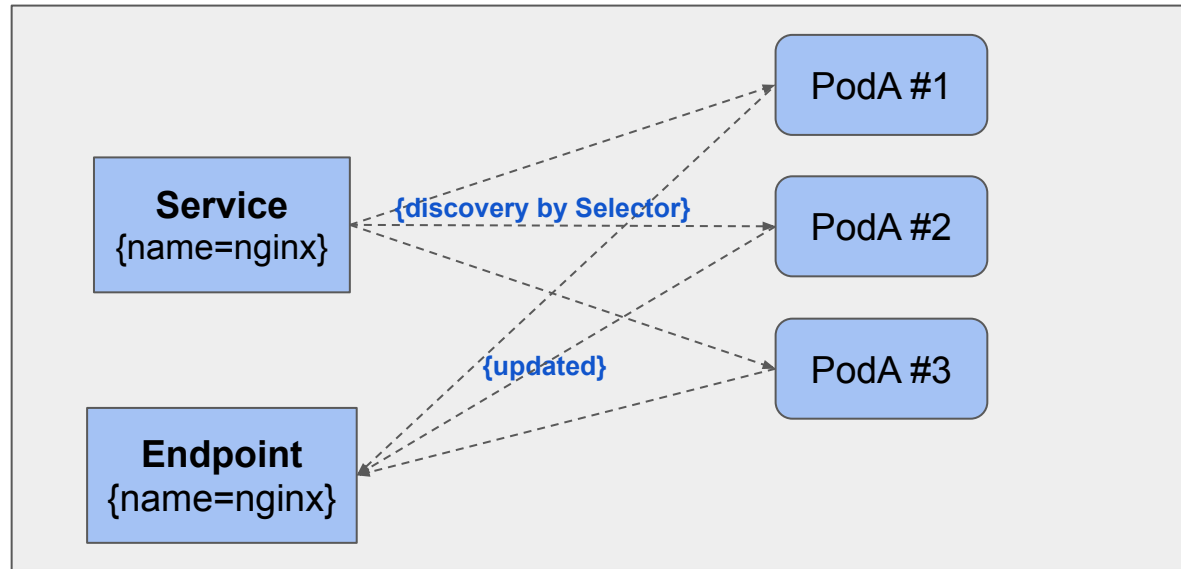
```
✓ kubectl describe svc nginx
Name:          nginx
Namespace:     default
Labels:        app=nginx
Annotations:   <none>
Selector:      app=nginx
Type:          ClusterIP
IP:            10.98.167.206
Port:          <unset> 8080/TCP
TargetPort:    80/TCP
Endpoints:     10.1.2.26:80
Session Affinity: None
Events:        <none>
```

```
✓ kubectl describe endpoints nginx
Name:          nginx
Namespace:     default
Labels:        app=nginx
Annotations:   endpoints.kubernetes.io/last-change-
Subsets:
  Addresses:          10.1.2.26
  NotReadyAddresses:  <none>
  Ports:
    Name      Port      Protocol
    ----      -
    <unset>   80       TCP
Events:        <none>
```

Service object named "**nginx**", which targets TCP port 80 on any Pod with the **app=nginx** label.

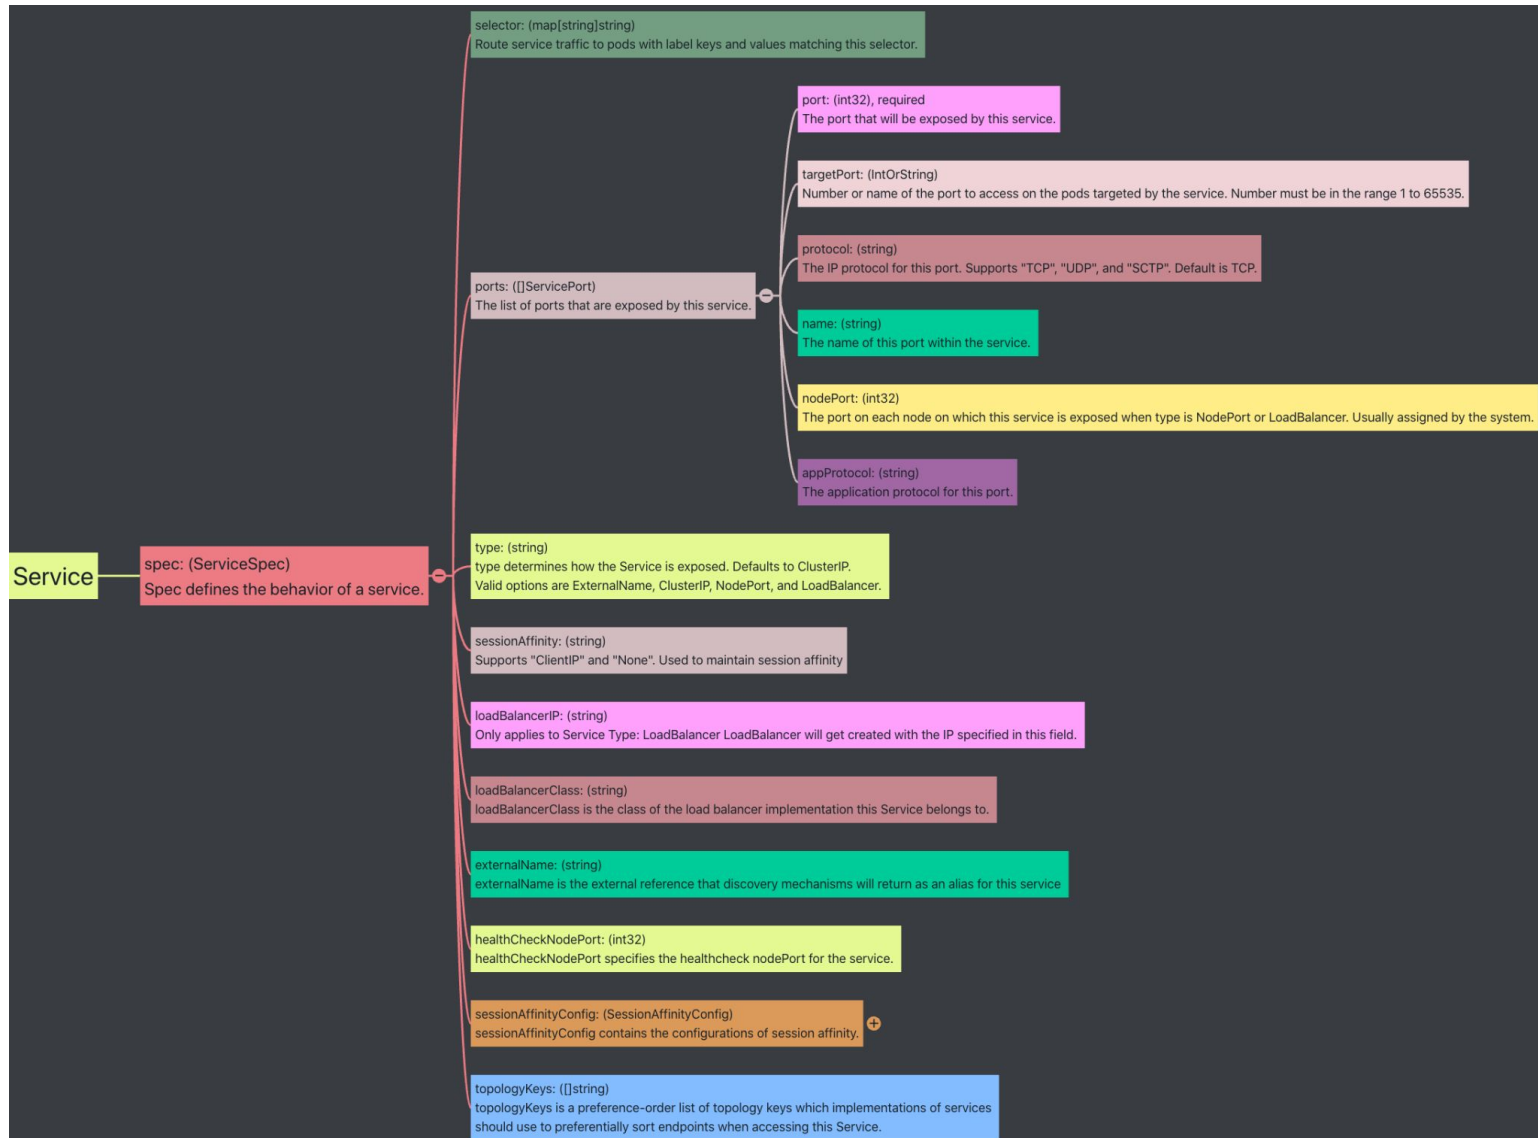
Notes: label app=nginx or run=nginx

Service and Endpoints

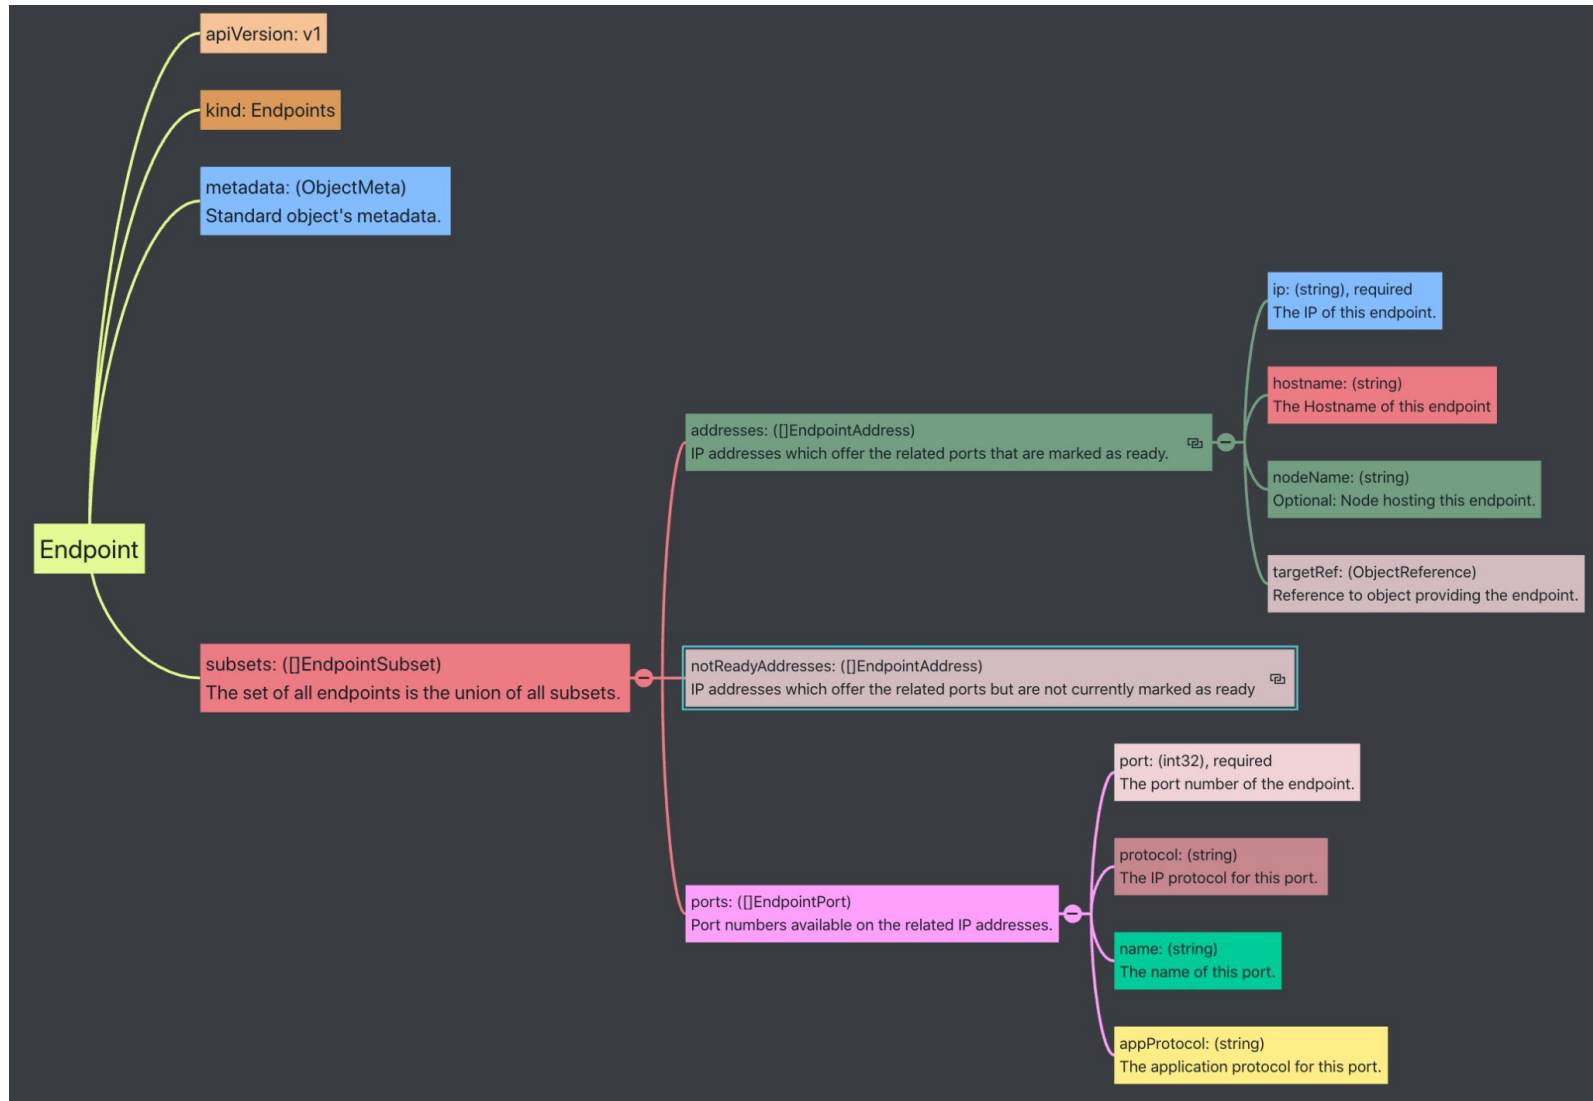


- The **controller** for the Service's selector will be evaluated continuously and the results will be POSTed to an Endpoints object.
- When a Pod dies, it is automatically removed from the endpoints.
- New Pods matching the Service's selector will automatically get added to the endpoints.
- Endpoints track the IP Addresses of the objects the service send traffic to.
- **Services without selectors**

Service Object (partial)



Endpoints Object



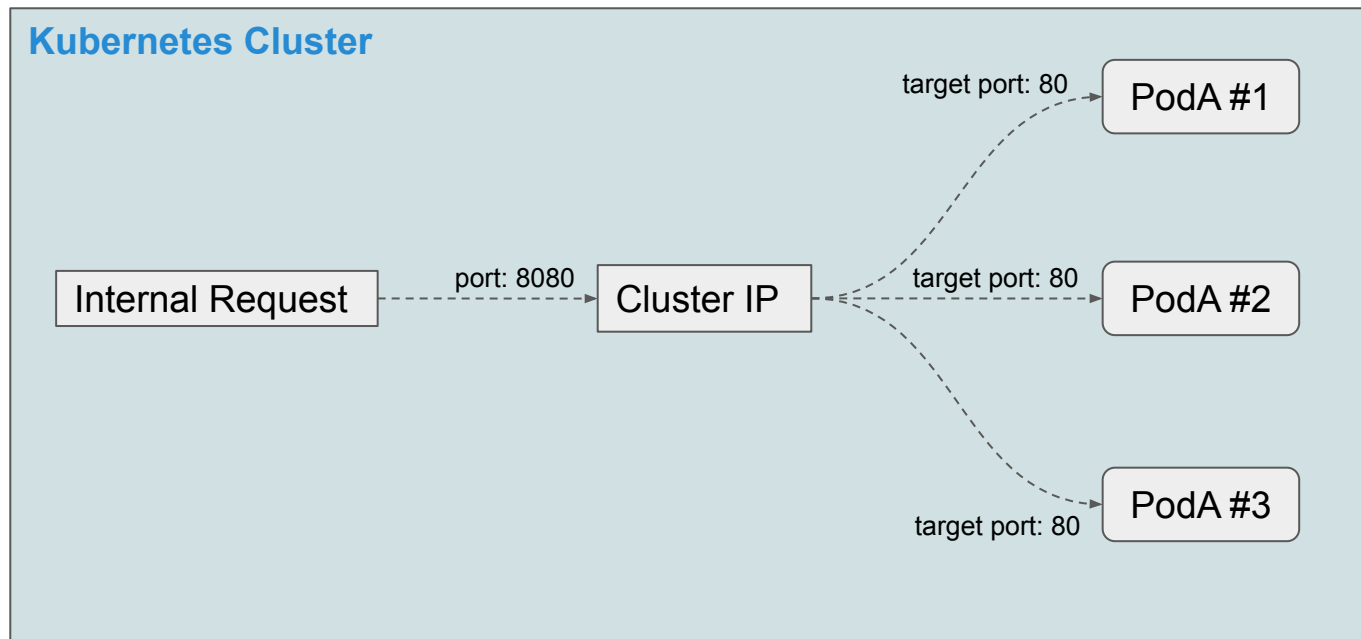
Publishing Services Using ServiceTypes

- Kubernetes [ServiceTypes](#) allow you to specify what kind of Service you want.

Types	Client	Summary
ClusterIP	Internal	This is the default ServiceType . Exposes the Service on a cluster-internal IP.
NodePort	External	Exposes the Service on each Node's IP at a static port .
LoadBalancer	External	Exposes the Service externally using a cloud provider's load balancer .
ExternalName	External	Maps the Service to the contents of the externalName field (e.g. foo.bar.example.com)

ServiceType - ClusterIP

- Exposes the Service on a cluster-internal IP.
- Choosing this type makes the Service **only reachable from within the cluster**.
- This is the **default ServiceType**.



ServiceType - ClusterIP Examples

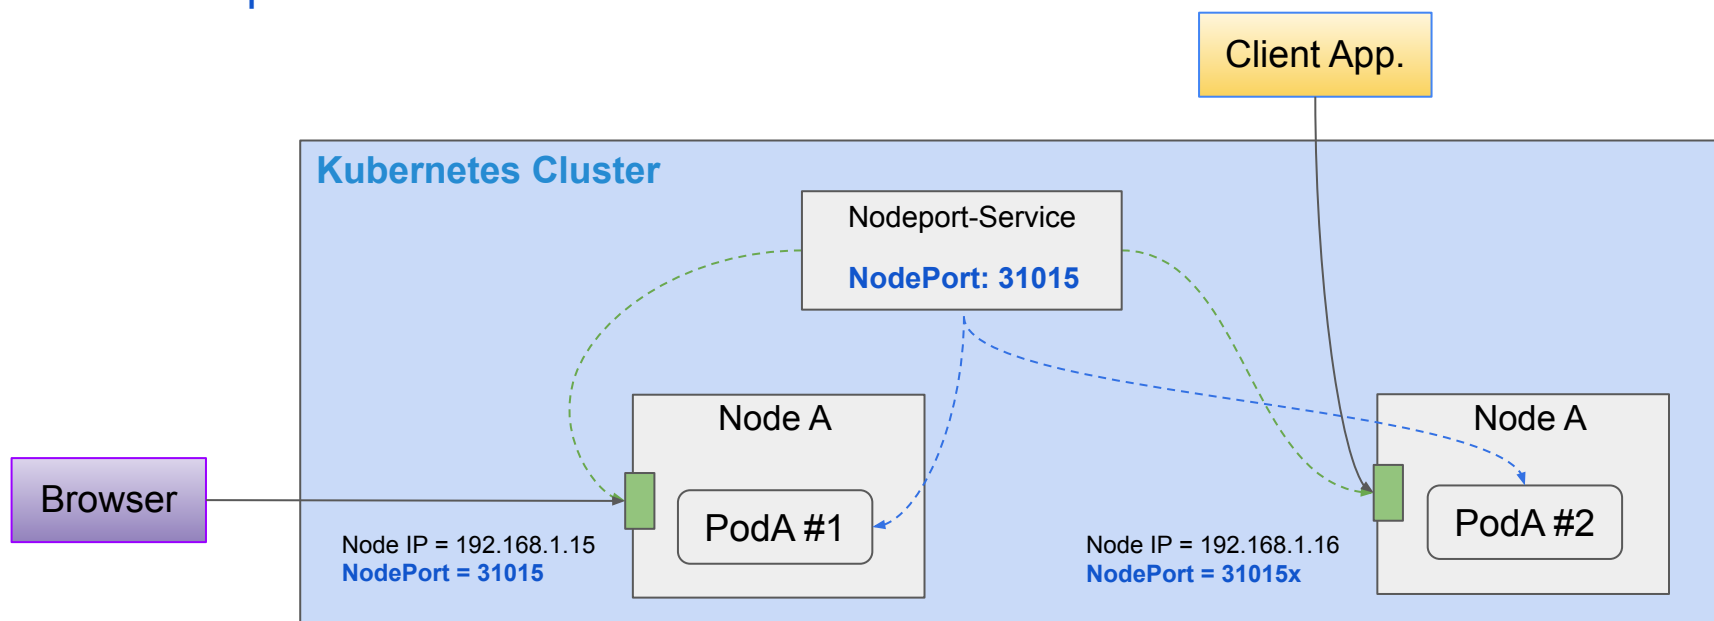
```
apiVersion: v1
kind: Service
metadata:
  labels:
    app: nginx
  name: nginx
  namespace: default
spec:
  ports:
  - port: 8080
    protocol: TCP
    targetPort: 80
  selector:
    app: nginx
  type: ClusterIP
```

nginx-service.yml

- `kubectl create deploy nginx --image=nginx`
- `kubectl expose deploy nginx --name nginx --port 8080 --target-port 80`
- Within Busybox
 - `nslookup nginx`
 - `curl nginx:8080`
- `kubectl scale --replicas=3 deploy nginx`
- `kubectl get po -lapp=nginx -owide`

ServiceType - NodePort

- Exposing Kubernetes services to external clients.
- Kubernetes exposes the Service on each Node's IP at a static port (the **NodePort** in the default range of **30000-32767**).
- You can specify node port or let Kubernetes automatic created.
- External clients can access the Service from outside the cluster by requesting **<NodeIp>:<NodePort>**



ServiceType - NodePort (Random Port)

```
apiVersion: v1
kind: Service
metadata:
  name: nodeport-service
  labels:
    app: nginx
spec:
  selector:
    app: nginx
  ports:
    - port: 8080
      targetPort: 80
  type: NodePort
```

nodeport-service.yml

```
✓ kubectl describe svc nodeport-service
Name: nodeport-service
Namespace: default
Labels: app=nginx
Annotations: kubectrl.kubernetes.io/last-applied-conf: {"apiVersion": "v1", "kind": "Service", "service", "namespace": "default"}...
Selector: app=nginx
Type: NodePort
IP: 10.96.31.93
LoadBalancer Ingress: localhost
Port: <unset> 8080/TCP
TargetPort: 80/TCP
NodePort: <unset> 30997/TCP
Endpoints: 10.1.2.38:80,10.1.2.39:80,10.1.2.40:80
Session Affinity: None
External Traffic Policy: Cluster
Events: <none>
```

```
✓ kubectl get po -l app=nginx -owide
NAME READY STATUS RESTARTS AGE IP
nginx-65f88748fd-djflr 1/1 Running 0 2d16h 10.1.2.40
nginx-65f88748fd-j9z4f 1/1 Running 0 2d16h 10.1.2.39
nginx-65f88748fd-rt2h4 1/1 Running 3 10d 10.1.2.38
```

```
✓ kubectl describe ep nodeport-service
Name: nodeport-service
Namespace: default
Labels: app=nginx
Annotations: endpoints.kubernetes.io/last-change-trigger-time: 2021-05-25T10:14:07Z
Subsets:
Addresses: 10.1.2.38,10.1.2.39,10.1.2.40 ← Pod ip(s)
NotReadyAddresses: <none>
Ports:
  Name Port Protocol
  ----
  <unset> 80 TCP
Events: <none>
```

- **curl localhost:30997**

ServiceType - NodePort (Specify port number)

```
apiVersion: v1
kind: Service
metadata:
  name: nodeport-service
  labels:
    app: nginx
spec:
  selector:
    app: nginx
  ports:
    - port: 8080
      targetPort: 80
      nodePort: 31001
  type: NodePort
```

nodeport31001-service.yml

```
✓ kubectl describe svc nodeport31001-service
Name: nodeport31001-service
Namespace: default
Labels: app=nginx
Annotations: kubectll.kubernetes.io/last-applied-conf
              {"apiVersion": "v1", "kind": "Service", "
1001-service", "namespace": "defa...
Selector: app=nginx
Type: NodePort
IP: 10.109.187.59
LoadBalancer Ingress: localhost
Port: <unset> 8080/TCP
TargetPort: 80/TCP
NodePort: <unset> 31001/TCP
Endpoints: 10.1.2.38:80,10.1.2.39:80,10.1.2.40:80
Session Affinity: None
External Traffic Policy: Cluster
Events: <none>
```

```
✓ kubectl get po -l app=nginx -owide
NAME                                READY   STATUS    RESTARTS   AGE   IP
nginx-65f88748fd-djflr             1/1     Running   0           2d16h 10.1.2.40
nginx-65f88748fd-j9z4f             1/1     Running   0           2d16h 10.1.2.39
nginx-65f88748fd-rt2h4             1/1     Running   3           10d    10.1.2.38
```

```
✓ kubectl describe ep nodeport31001-service
Name: nodeport31001-service
Namespace: default
Labels: app=nginx
Annotations: endpoints.kubernetes.io/last-change-trigger-time: 2021-05-25T10:49:03Z
Subsets:
Addresses: 10.1.2.38,10.1.2.39,10.1.2.40 ← Pod ip(s)
NotReadyAddresses: <none>
Ports:
  Name      Port      Protocol
  ----      -
  <unset>   80        TCP
Events: <none>
```

Warning - Need to take care of possible port collisions yourself.

Create NodePort Service Using Imperative Way

- `kubectl create svc nodeport --node-port=31013 --tcp=8080:80 nginx`

```
✓ kubectl describe svc nginx
Name: nginx
Namespace: default
Labels: app=nginx
Annotations: <none>
Selector: app=nginx
Type: NodePort
IP: 10.101.202.91
LoadBalancer Ingress: localhost
Port: 8080-80 8080/TCP
TargetPort: 80/TCP
NodePort: 8080-80 31013/TCP
Endpoints: 10.1.2.38:80,10.1.2.39:80,10.1.2.40:80
Session Affinity: None
External Traffic Policy: Cluster
Events: <none>
```

Note: "nginx" was used as a Selector.

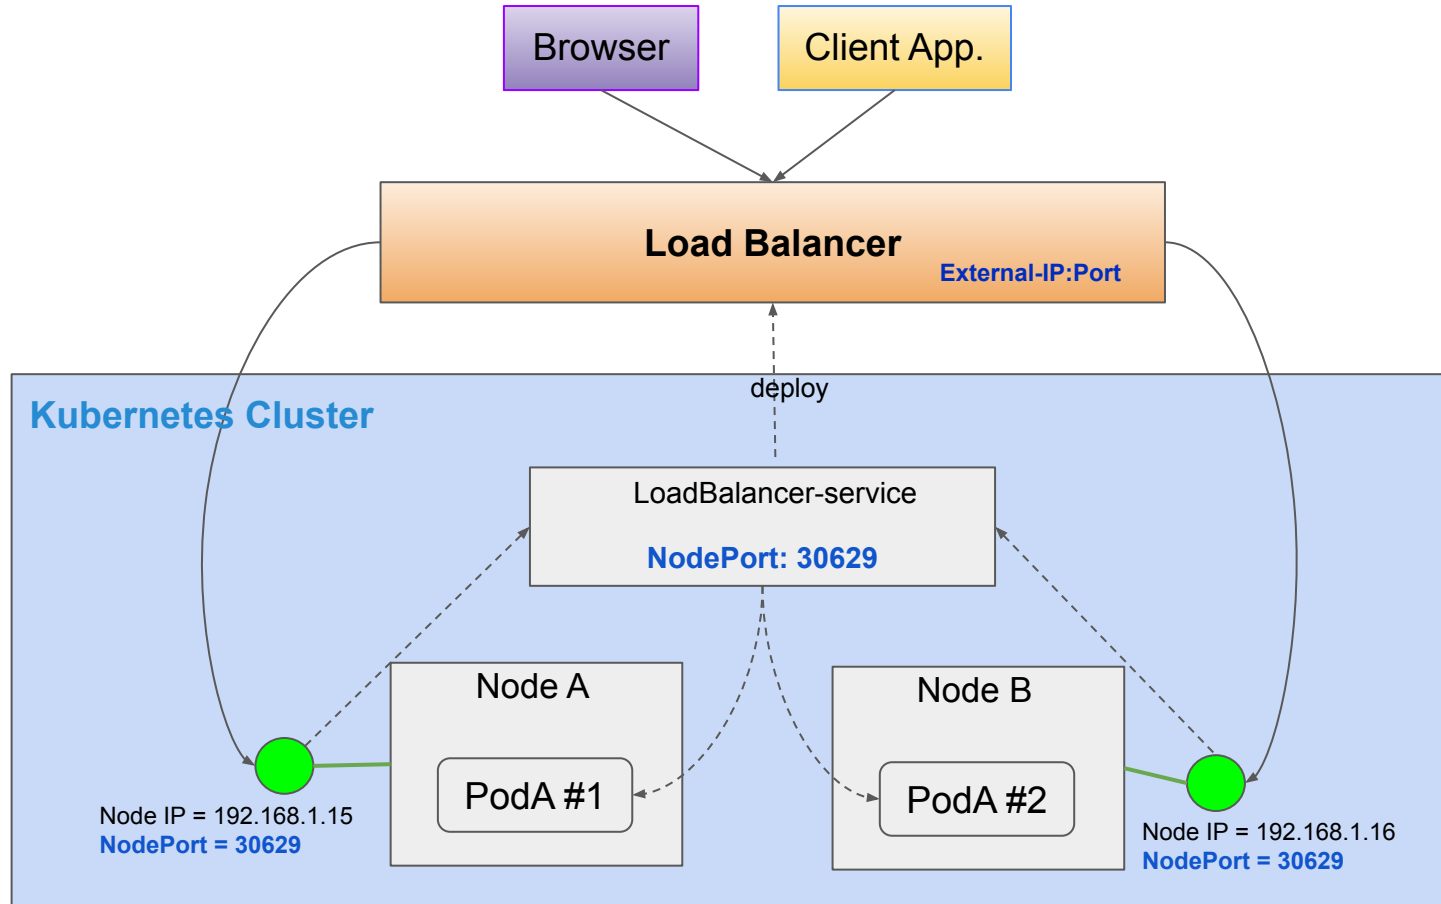
- `kubectl expose deploy nginx --port=8080 --target-port=80 --type NodePort --name nginx-svc2`

```
✓ kubectl describe svc nginx-svc2
Name: nginx-svc2
Namespace: default
Labels: app=nginx
Annotations: <none>
Selector: app=nginx
Type: NodePort
IP: 10.104.233.73
LoadBalancer Ingress: localhost
Port: <unset> 8080/TCP
TargetPort: 80/TCP
NodePort: <unset> 32652/TCP
Endpoints: 10.1.2.38:80,10.1.2.39:80,10.1.2.40:80
Session Affinity: None
External Traffic Policy: Cluster
Events: <none>
```


ServiceType - LoadBalancer

- A LoadBalancer service is the [standard way to expose a service to the internet](#).
- On cloud providers, setting the type field to “[LoadBalancer](#)” automatically [deploys an external load balancer](#) and [provided a load balancer](#) for your Service.
- The exact implementation of a LoadBalancer is dependent on your cloud provider.
- The load balancer will have its own unique, publicly accessible IP address.
- Client access your services through the load balancer’s IP address.
- a LoadBalancer service is an extension of a NodePort service.
- [metadata.annotations](#) are used to define the properties and features of each cloud provider.

ServiceType - LoadBalancer



<https://kubernetes.io/docs/concepts/services-networking/service/#internal-load-balancer>

<https://kubernetes.io/docs/tasks/access-application-cluster/create-external-load-balancer/#preserving-the-client-source-ip>

ServiceType - LoadBalancer Local Example

```
apiVersion: v1
kind: Service
metadata:
  name: lb-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 80
  type: LoadBalancer
```

lb-service.yml

```
✓ kubectl describe svc lb-service
Name: lb-service
Namespace: default
Labels: <none>
Annotations: kubectl.kubernetes.io/last-applied-configuration: {"apiVersion": "v1", "kind": "Service", "metadata": {"name": "lb-service", "namespace": "default"}, "spec": {"ports": [{"port": 8080, "protocol": "TCP", "targetPort": 80}], "selector": {"app": "nginx"}, "type": "LoadBalancer"}}
Selector: app=nginx
Type: LoadBalancer
IP: 10.109.242.6
LoadBalancer Ingress: localhost
Port: <unset> 8080/TCP
TargetPort: 80/TCP
NodePort: <unset> 30629/TCP
Endpoints: 10.1.2.38:80,10.1.2.39:80,10.1.2.40:80
Session Affinity: None
External Traffic Policy: Cluster
Events: <none>
```

```
✓ kubectl get ep lb-service
NAME ENDPOINTS AGE
lb-service 10.1.2.38:80,10.1.2.39:80,10.1.2.40:80 14m
```

- curl localhost:8080

Accessing LoadBalancer Service

```
✓ kubectl get svc nodeport-service lb-service
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
nodeport-service	NodePort	10.106.1.65	<none>	8080:30567/TCP	6h
lb-service	LoadBalancer	10.109.242.6	localhost	8080:30629/TCP	31m

```
✓ kubectl get ep nodeport-service lb-service
```

NAME	ENDPOINTS	AGE
nodeport-service	10.1.2.38:80,10.1.2.39:80,10.1.2.40:80	6h2m
lb-service	10.1.2.38:80,10.1.2.39:80,10.1.2.40:80	33m

```
✓ curl localhost:8080 !
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
```

LoadBalancer

```
✓ curl localhost:30629 ! LoadBalancer
```

```
curl: (7) Failed to connect to localhost port 30629: Connection refused
```

```
✓ curl localhost:30567 ! NodePort
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
```

ServiceType - LoadBalancer Service YAML

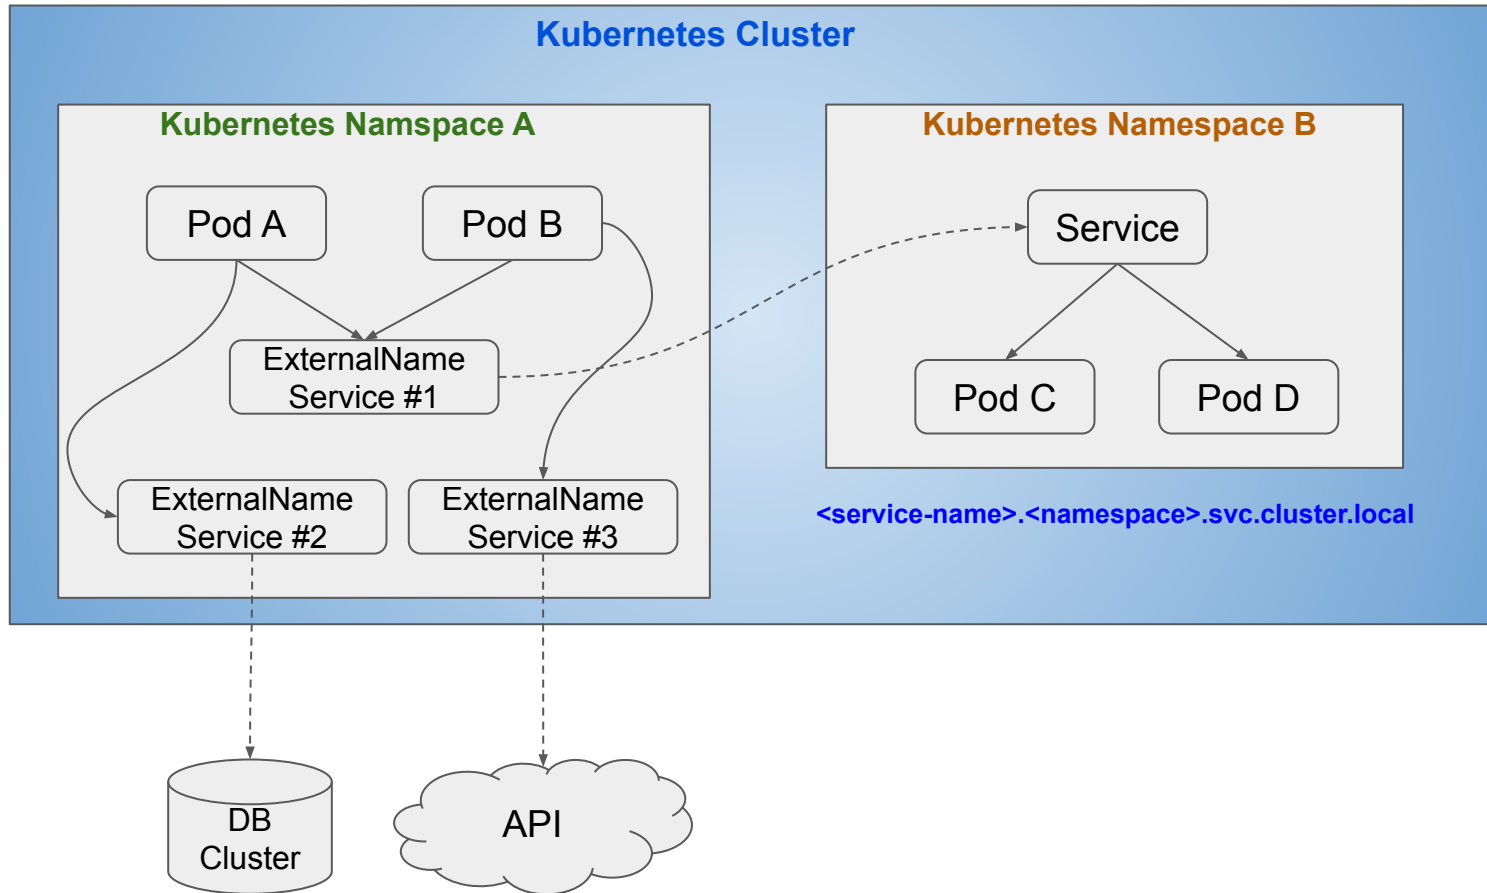
```
apiVersion: v1
kind: Service
metadata:
  annotations:
    kubectl.kubernetes.io/last-applied-configuration: |

{"apiVersion":"v1","kind":"Service","metadata":{"annotations":{},"name":"lb-service","namespace":"default"},"spec":{"ports":[{"port":8080,"protocol":"TCP","targetPort":80}],"selector":{"app":"nginx"},"type":"LoadBalancer"}}
  creationTimestamp: "2021-05-27T03:07:28Z"
  name: lb-service
  namespace: default
  resourceVersion: "3046010"
  selfLink: /api/v1/namespaces/default/services/lb-service
  uid: ab52795a-be98-11eb-90fc-025000000001
spec:
  clusterIP: 10.107.160.172
  externalTrafficPolicy: Cluster
  ports:
    - nodePort: 31033
      port: 8080
      protocol: TCP
      targetPort: 80
  selector:
    app: nginx
  sessionAffinity: None
  type: LoadBalancer
status:
  loadBalancer:
    ingress:
      - hostname: localhost
```

ServiceType - ExternalName

- Required to refer to [external services](#) from applications inside your Kubernetes cluster.
- External services may be,
 - [Services in another Kubernetes Namespace.](#)
 - [Services outside a Kubernetes Cluster.](#)
- An ExternalName Service is a special case of Service that does not have selectors and uses DNS names instead.
- Services of type ExternalName [map a Service to a DNS name \(not an IP/port\).](#)

ServiceType - ExternalName



ServiceType - ExternalName Example #1

- Call other namespace service.

```
apiVersion: v1
kind: Service
metadata:
  name: other-ns-service
spec:
  type: ExternalName
  externalName: hello.test.svc.cluster.local
```

other-ns-service.yml

- `kubectl create ns test`
- `kubectl run hello --image=gcr.io/google-samples/hello-app:1.0 -n test`
- `kubectl scale --replicas=3 deploy hello -n test`
- `kubectl expose deploy hello --port=8090 --target-port=8080 --name hello -n test`
- `kubectl apply -f other-ns-service.yml`
- `curl other-ns-service:8090 # within busybox`

ServiceType - ExternalName Example #2

- Call external API

```
apiVersion: v1
kind: Service
metadata:
  name: testapi
spec:
  type: ExternalName
  externalName: postman-echo.com
```

postman-echo-service.yml

```
import requests
response =
requests.get('http://testapi/get?arg1=value123')
print(response.text)
```

client.py

```
FROM python:3.9-slim
WORKDIR /app
COPY requirements.txt ./
RUN pip install --no-cache-dir -r requirements.txt
COPY . .
CMD ["python", "./client.py"]
```

Dockerfile

- `kubectl apply -f postman-echo-service.yml`
- `curl testapi/get?a=b` #within busybox
- `kubectl run apiclient --restart=Never --image=apiclient:v1.0` # check pod's log

Services without selectors - Mapping a hostname to an IP

```
kind: "Service"
apiVersion: "v1"
metadata:
  name: "mariadb"
spec:
  ports:
    - name: "ext-mariadb"
  protocol: "TCP"
  port: 3333
  targetPort: 3306
```

external-mariadb.yml

```
kind: "Endpoints"
apiVersion: "v1"
metadata:
  name: "mariadb"
subsets:
  - addresses:
      - ip: "192.168.1.16"
    ports:
      - port: 3306
        name: "ext-mariadb"
```

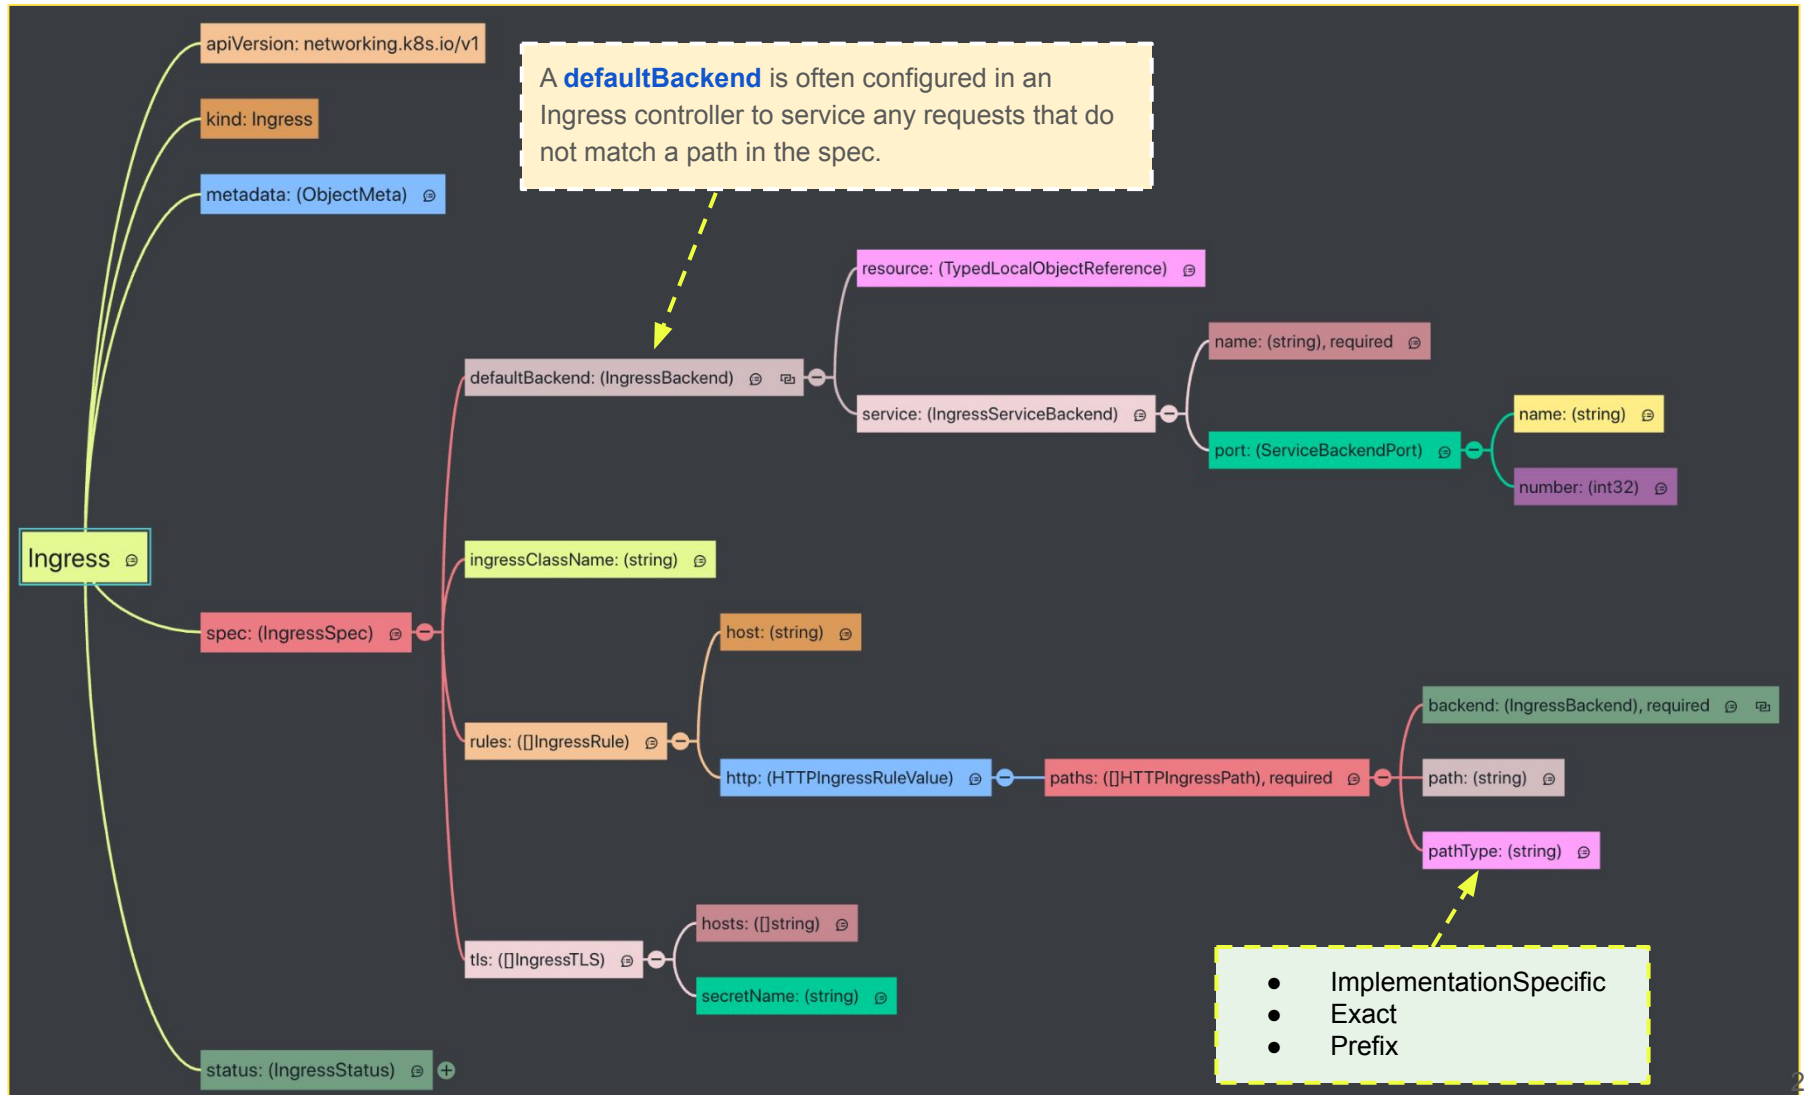
```
/ # telnet -l root mariadb 3333
X
5.5.5-10.5.8-MariaDBW;`R[Ri800-y|"0kKu,w6G@mysql_native_password
Connection closed by foreign host
```

- Because this Service has **no selector**, the corresponding Endpoints object is not created automatically.

Ingress Overview

- Ingress is an Kubernetes [API object](#) that exposes HTTP and HTTPS routes from outside the cluster to services within the cluster.
- Ingress is not a Service type, but it a [reverse proxy](#) and [single entry-point](#) to your cluster that routes the request to different services.
- it can expose multiple services under the same IP address.
- Traffic routing is controlled by [rules](#) defined on the [Ingress resource](#).
- An Ingress may be configured to give Services externally-reachable URLs, [load balance traffic](#), [terminate SSL / TLS](#), and offer [name-based virtual hosting](#).
- You must have an [Ingress controller](#) to satisfy an Ingress. Only creating an Ingress resource has no effect.
- Ingress frequently uses [annotations](#) to configure some options depending on the Ingress controller.

Ingress Object



Hostname and PathType

Hosts can be precise matches (for example “foo.bar.com”) or a wildcard (for example “*.foo.com”).

Host	Host header	Match?
*.foo.com	bar.foo.com	Matches based on shared suffix
*.foo.com	baz.bar.foo.com	No match, wildcard only covers a single DNS label
*.foo.com	foo.com	No match, wildcard only covers a single DNS label

Kind	Path(s)	Request path(s)	Matches?
Prefix	/	(all paths)	Yes
Exact	/foo	/foo	Yes
Exact	/foo	/bar	No
Exact	/foo	/foo/	No
Exact	/foo/	/foo	No
Prefix	/foo	/foo , /foo/	Yes
Prefix	/foo/	/foo , /foo/	Yes
Prefix	/aaa/bb	/aaa/bbb	No
Prefix	/aaa/bbb	/aaa/bbb	Yes
Prefix	/aaa/bbb/	/aaa/bbb	Yes, ignores trailing slash
Prefix	/aaa/bbb	/aaa/bbb/	Yes, matches trailing slash
Prefix	/aaa/bbb	/aaa/bbb/ccc	Yes, matches subpath
Prefix	/aaa/bbb	/aaa/bbbxyz	No, does not match string prefix
Prefix	/ , /aaa	/aaa/ccc	Yes, matches /aaa prefix
Prefix	/ , /aaa , /aaa/bbb	/aaa/bbb	Yes, matches /aaa/bbb prefix
Prefix	/ , /aaa , /aaa/bbb	/ccc	Yes, matches / prefix
Prefix	/aaa	/ccc	No, uses default backend
Mixed	/foo (Prefix), /foo (Exact)	/foo	Yes, prefers Exact

Ingress Controller

- An ingress controller is responsible for reading the [Ingress Resource](#) information and processing that data accordingly.
- A collection of [routing rules](#) that govern [how external users access services](#) running in a Kubernetes cluster.
- Different ingress controllers have extended the specification in different ways to support additional use cases.
- Note that an ingress controller typically [doesn't eliminate](#) the need for an external load balancer — the ingress controller simply adds an additional layer of routing and control behind the load balancer.
- An Ingress controller is bootstrapped with some load balancing policy settings that it applies to all Ingress.
- The most basic Ingress is the [NGINX Ingress Controller](#).

Nginx Ingress Controller

- Install for Kubernetes 1.14
 - `kubectl apply -f`
<https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v0.35.0/deploy/static/provider/cloud/deploy.yaml>
- `kubectl apply -f`
<https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v0.47.0/deploy/static/provider/cloud/deploy.yaml>
- Ingress frequently uses annotations to configure some options depending on the Ingress controller.
- Different Ingress controller support different annotations.

```
✔ kubectl get all -n ingress-nginx
```

NAME	READY	STATUS	RESTARTS	AGE
pod/ingress-nginx-admission-create-vrxdw	0/1	Completed	0	3d23h
pod/ingress-nginx-admission-patch-5z2jb	0/1	Completed	0	3d23h
pod/ingress-nginx-controller-68556b9795-gj4n5	1/1	Running	2	3d23h

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/ingress-nginx-controller	LoadBalancer	10.96.31.105	localhost	80:31617/TCP,443:31755/TCP	3d23h
service/ingress-nginx-controller-admission	ClusterIP	10.98.31.172	<none>	443/TCP	3d23h

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/ingress-nginx-controller	1/1	1	1	3d23h

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/ingress-nginx-controller-68556b9795	1	1	1	3d23h

NAME	COMPLETIONS	DURATION	AGE
job.batch/ingress-nginx-admission-create	1/1	11s	3d23h
job.batch/ingress-nginx-admission-patch	1/1	16s	3d23h

Additional controllers

- [AKS Application Gateway Ingress Controller](#) is an ingress controller that configures the [Azure Application Gateway](#).
- [Ambassador](#) API Gateway is an [Envoy](#)-based ingress controller.
- [Apache APISIX ingress controller](#) is an [Apache APISIX](#)-based ingress controller.
- [Avi Kubernetes Operator](#) provides L4-L7 load-balancing using [VMware NSX Advanced Load Balancer](#).
- The [Citrix ingress controller](#) works with Citrix Application Delivery Controller.
- [Contour](#) is an [Envoy](#) based ingress controller.
- [EnRoute](#) is an [Envoy](#) based API gateway that can run as an ingress controller.
- [F5 BIG-IP Container Ingress Services for Kubernetes](#) lets you use an Ingress to configure F5 BIG-IP virtual servers.
- [Gloo](#) is an open-source ingress controller based on [Envoy](#), which offers API gateway functionality.
- [HAProxy Ingress](#) is an ingress controller for [HAProxy](#).
- The [HAProxy Ingress Controller for Kubernetes](#) is also an ingress controller for [HAProxy](#).
- [Istio Ingress](#) is an [Istio](#) based ingress controller.
- The [Kong Ingress Controller for Kubernetes](#) is an ingress controller driving [Kong Gateway](#).
- The [NGINX Ingress Controller for Kubernetes](#) works with the [NGINX](#) webserver (as a proxy).
- [Skipper](#) HTTP router and reverse proxy for service composition, including use cases like Kubernetes Ingress, designed as a library to build your custom proxy.
- The [Traefik Kubernetes Ingress provider](#) is an ingress controller for the [Traefik](#) proxy.
- [Tyk Operator](#) extends Ingress with Custom Resources to bring API Management capabilities to Ingress. Tyk Operator works with the Open Source Tyk Gateway & Tyk Cloud control plane.
- [Voyager](#) is an ingress controller for [HAProxy](#).

Ingress - Example#1 No Host

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: ingress-nohost
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
    - http:
        paths:
          - path: /hello
            backend:
              serviceName: hello
              servicePort: 80
          - path: /web
            backend:
              serviceName: nodeport-service
              servicePort: 8080
```

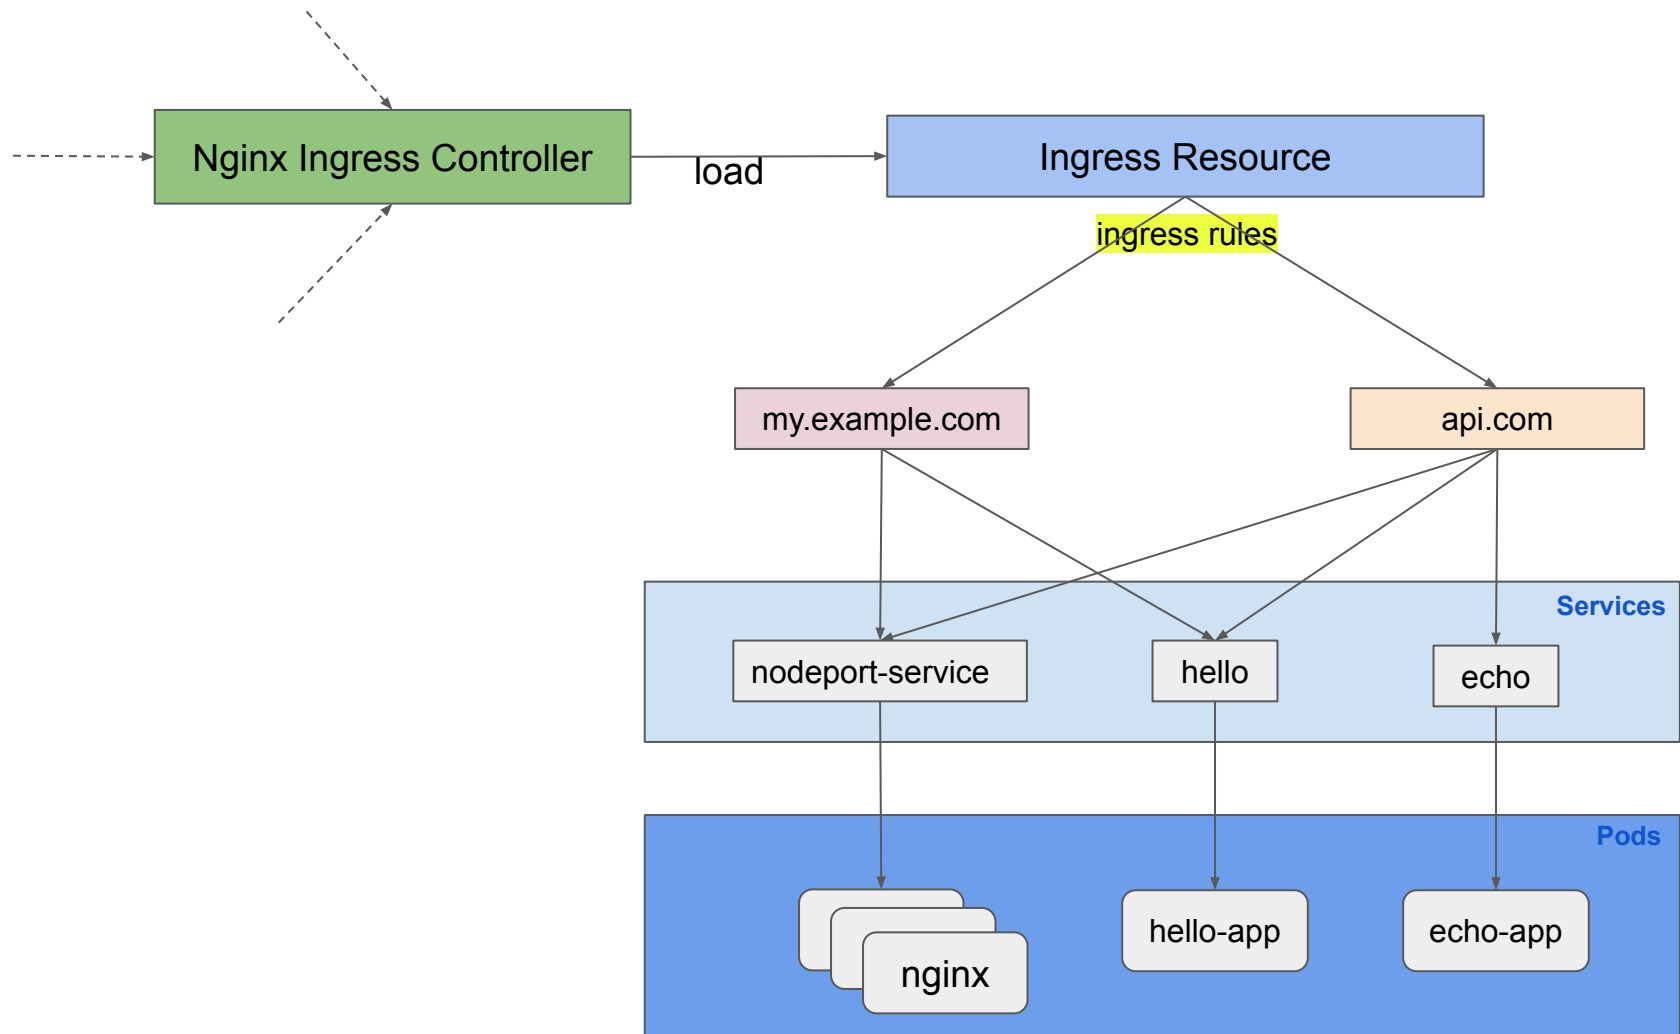
ingress-nohost.yml

- this rule applies to all inbound HTTP traffic through the IP address specified.

Ingress - Example#1 No Host Steps

- `kubectl run hello --image=gcr.io/google-samples/hello-app:1.0`
- `kubectl scale --replicas=3 deploy hello`
- `kubectl expose deploy hello --port=80 --target-port=8080 --name hello`
- `kubectl apply -f ingress-nohost.yml`
- `curl localhost/hello`
- `kubectl delete deploy hello`
- `kubectl delete svc hello`
- `kubectl delete ingress ingress-nohost`

Ingress - Example#2 Name Based Virtual Hosting



Ingress - Example #2 Ingress Resource

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: ingress-multiple
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - host: my.example.com
    http:
      paths:
      - path: /nginx
        backend:
          serviceName: nodeport-service
          servicePort: 8080
      - path: /hello
        backend:
          serviceName: hello
          servicePort: 80
```

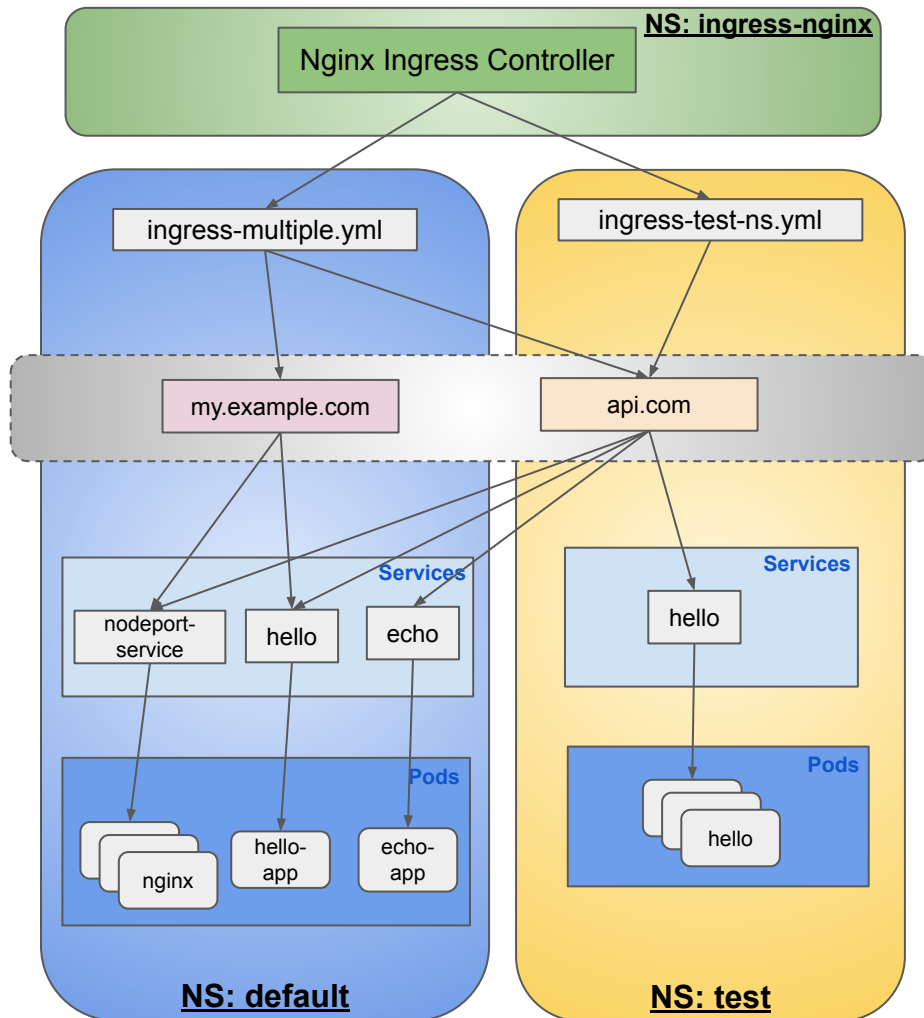
ingress-multiple.yml

```
- host: api.com
  http:
    paths:
    - path: /doc
      backend:
        serviceName: nodeport-service
        servicePort: 8080
    - path: /ping
      backend:
        serviceName: hello
        servicePort: 80
    - path: /echo
      backend:
        serviceName: echo
        servicePort: 80
```

Ingress - Example #2 Steps

- `kubectl create deployment hello-app --image=gcr.io/google-samples/hello-app:1.0`
- `kubectl create deployment echo-app --image=k8s.gcr.io/echoserver:1.4`
- `kubectl expose deployment hello-app --port=80 --target-port=8080 --name=hello`
- `kubectl expose deployment echo-app --port=80 --target-port=8080 --name=echo`
- `kubectl apply -f nodeport-service.yml`
- `kubectl logs -f -n ingress-nginx ingress-nginx-controller-68556b9795-gj4n5`
- `kubectl apply -f ingress-multiple.yml`
- `kubectl get ing`
- Add “127.0.0.1 my.example.com api.com” to /etc/hosts file
- Tesing
 - `curl --header "Host: my.example.com" http://localhost/hello`
 - `curl http://my.example.com/nginx`
 - `curl --header "Host: api.com" http://localhost/doc`
 - `curl --header "Host: api.com" http://localhost/echo?a=b`
 - `curl http://api.com/ping`

Ingress - Multiple Namespaces Ingress Resource



```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: ingress-test-ns
  namespace: test
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - host: my.example.com
    http:
      paths:
      - path: /hello-test
        backend:
          serviceName: hello
          servicePort: 8090
```

ingress-test-ns.yml

- `kubectl create deployment hello --image=gcr.io/google-samples/hello-app:1.0 -n test`
- `kubectl expose deploy hello --port=8090 --target-port=8080 --name hello -n test`
- `kubectl scale --replicas=3 deploy hello -n test`
- `curl my.example.com/hello` # service in default namespace
- `curl my.example.com/hello-test` # service in test namespace

End.